

Deep learning

Introduction

Hamid Beigy

Sharif university of technology

September 30, 2019

Table of contents

1 Course Information

2 Introduction

3 Success stories

4 Outline of course



Table of contents

1 Course Information

2 Introduction

3 Success stories

4 Outline of course



Course Information

- 1 Course name : Deep learning
- 2 The objective of **deep learning** is moving **Machine Learning** closer to one of its original goals: **Artificial Intelligence**.
- 3 Instructor : Hamid Beigy
Email : beigy@sharif.edu
- 4 Course Website:
<http://ce.sharif.edu/courses/98-99/1/ce718-1/>
- 5 Lectures: Sat-Mon (10:30-12:00)
- 6 TAs :

Fariba Lotfi	Email: fлотfi@ce.sharif.edu
Farzad Beizaee	Email:



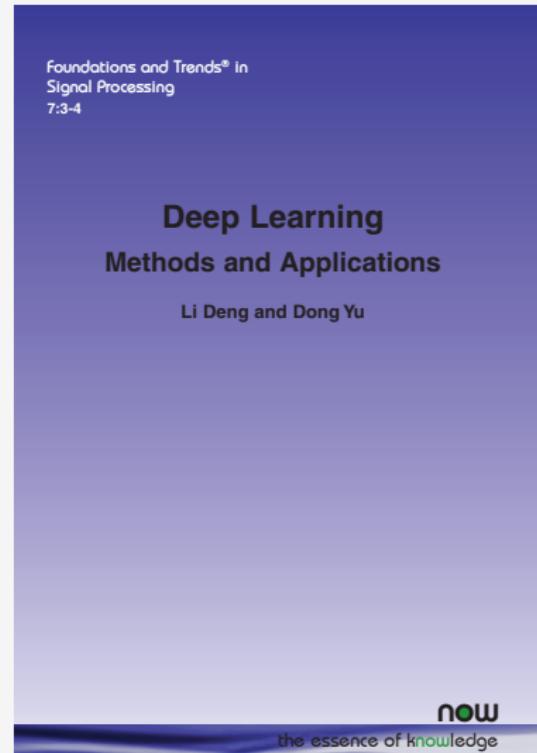
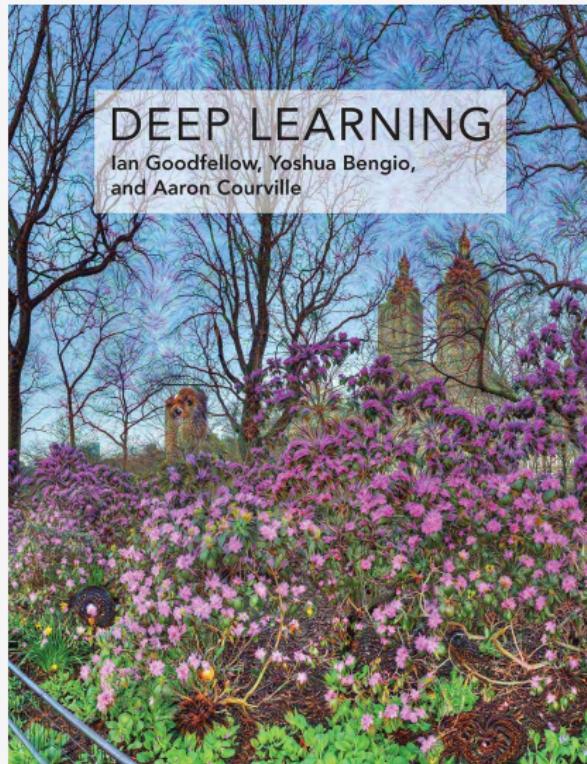
Course evaluation

■ Evaluation:

Mid-term exam	30%	1397/8/11
Final exam	30%	
Practical Assignments	30%	
Quiz	10%	
Paper	10%	



Main reference





References |

-  DENG, L., AND YU, D.
Deep learning: Methods and applications.
Foundations and Trends in Signal Processing 7, 3–4 (2013), 197–387.
-  GOODFELLOW, I., BENGIO, Y., AND COURVILLE, A.
Deep Learning.
MIT Press, 2016.



Relevant journals |

- 1 IEEE Trans on Pattern Analysis and Machine Intelligence
- 2 Journal of Machine Learning Research
- 3 Pattern Recognition
- 4 Machine Learning
- 5 Neural Networks
- 6 Neural Computation
- 7 Neurocomputing
- 8 IEEE Trans. on Neural Networks and Learning Systems
- 9 Annals of Statistics
- 10 Journal of the American Statistical Association
- 11 Pattern Recognition Letters
- 12 Artificial Intelligence



Relevant journals II

- [13 Data Mining and Knowledge Discovery](#)
- [14 IEEE Transaction on Cybernetics \(SMC-B\)](#)
- [15 IEEE Transaction on Knowledge and Data Engineering](#)
- [16 Knowledge and Information Systems](#)



Relevant conferences

- 1 Neural Information Processing Systems (NIPS)
- 2 International Conference on Machine Learning (ICML)
- 3 European Conference on Machine Learning (ECML)
- 4 Asian Conference on Machine Learning (ACML)
- 5 Conference on Learning Theory (COLT)
- 6 Algorithmic Learning Theory (ALT)
- 7 Conference on Uncertainty in Artificial Intelligence (UAI)
- 8 Practice of Knowledge Discovery in Databases (PKDD)
- 9 International Joint Conference on Artificial Intelligence (IJCAI)
- 10 IEEE International Conference on Data Mining series (ICDM)



Relevant packages and datasets

1 Packages:

- Keras <https://keras.io>
- TensorFlow <http://www.tensorflow.org/>
- Caffe <http://caffe.berkeleyvision.org>
- PyTorch <https://pytorch.org>

2 Datasets:

- UCI Machine Learning Repository
<http://archive.ics.uci.edu/ml/>
- MNIST: handwritten digits <http://yann.lecun.com/exdb/mnist/>
- 20 newsgroups <http://qwone.com/~jason/20Newsgroups/>



Table of contents

1 Course Information

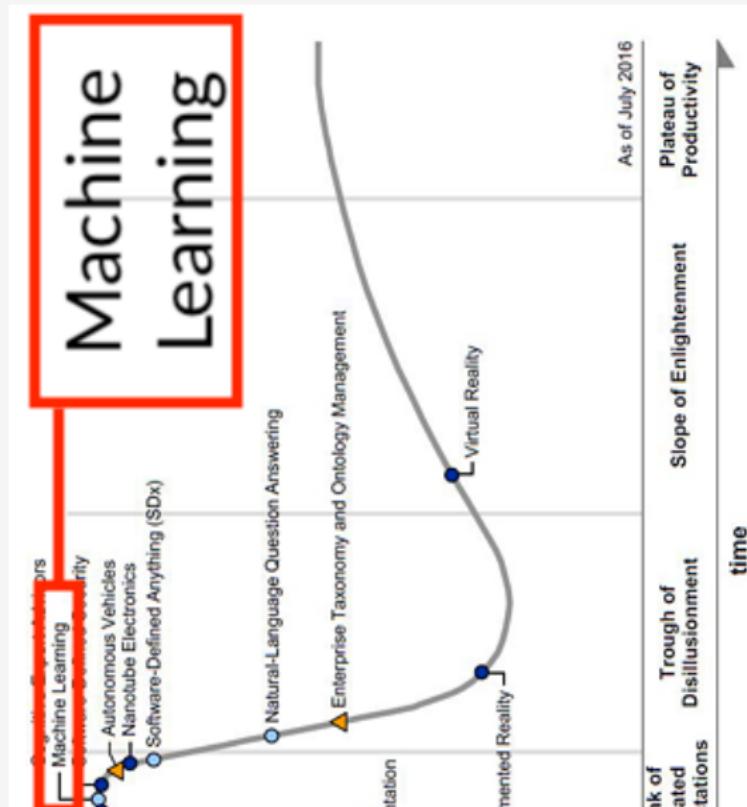
2 Introduction

3 Success stories

4 Outline of course

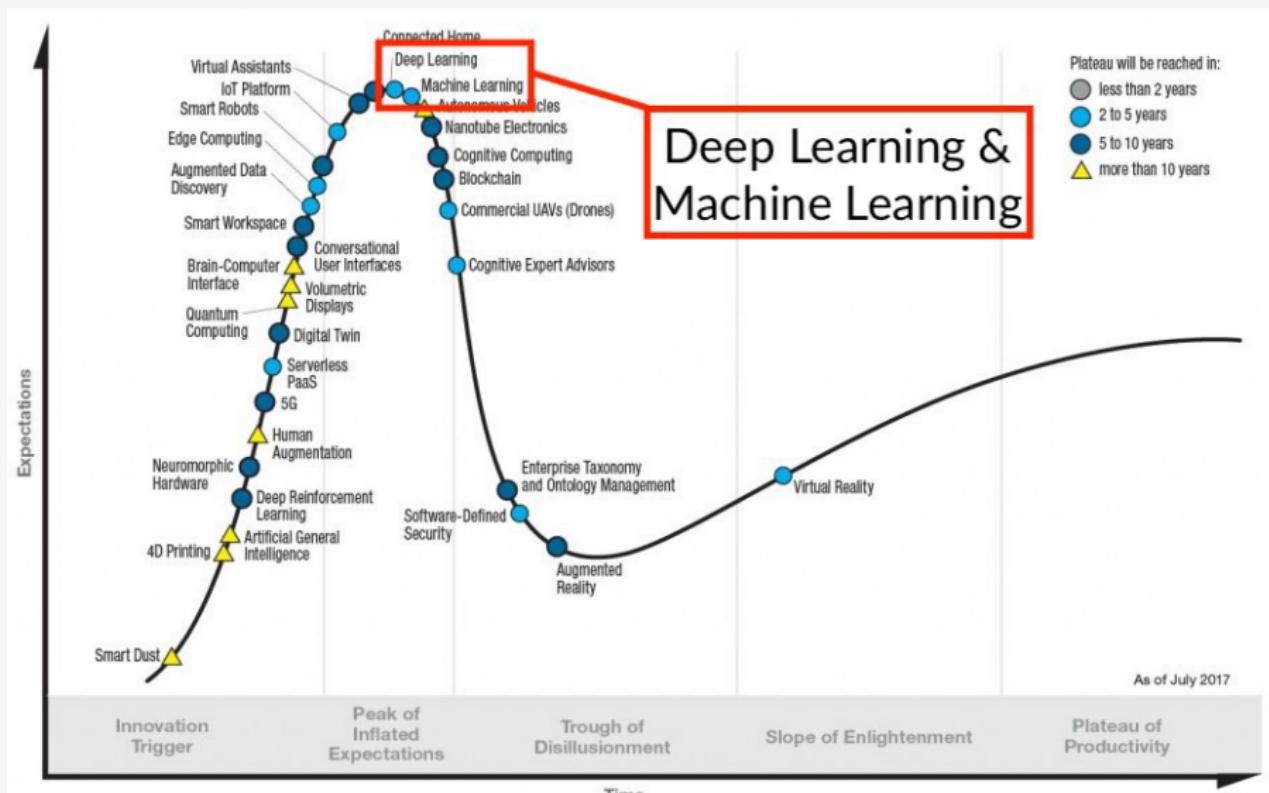


Gartner Hyper-Cycle of Emerging Technologies (2016)



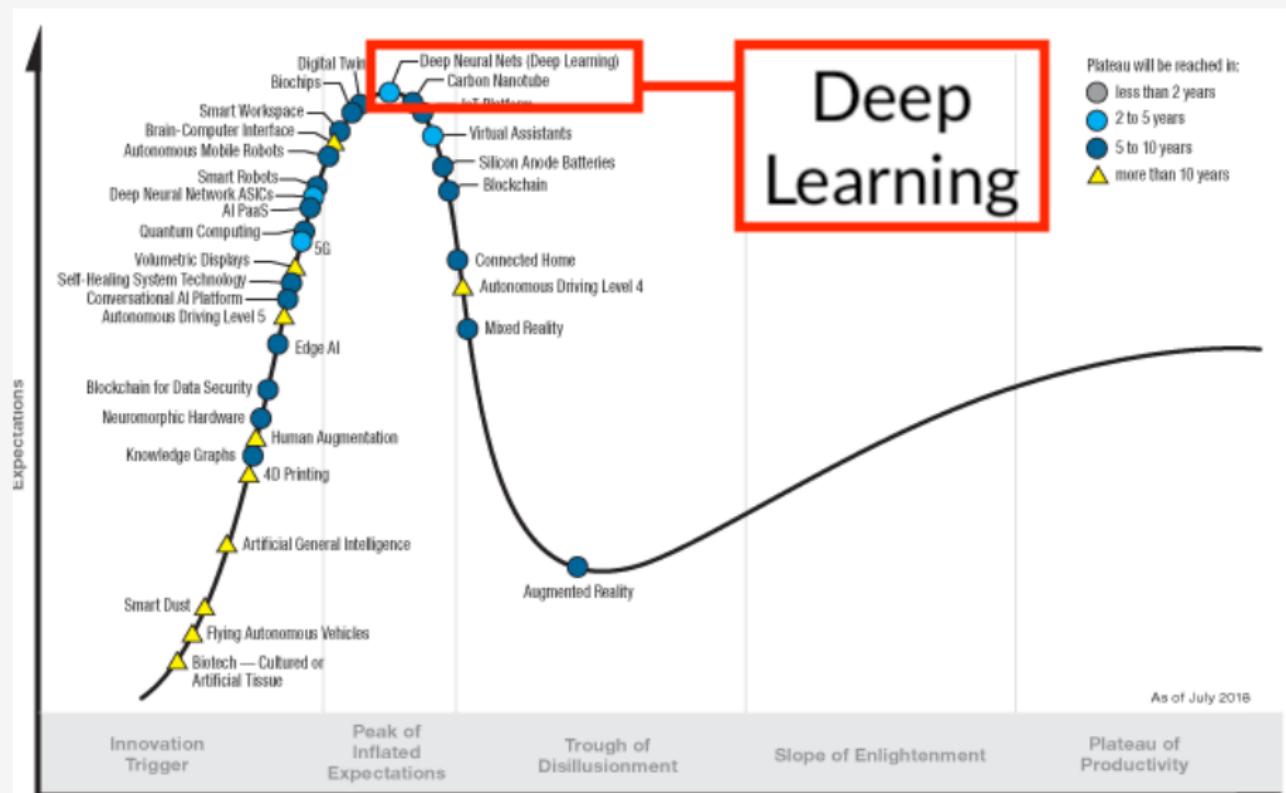


Gartner Hyper-Cycle of Emerging Technologies (2017)



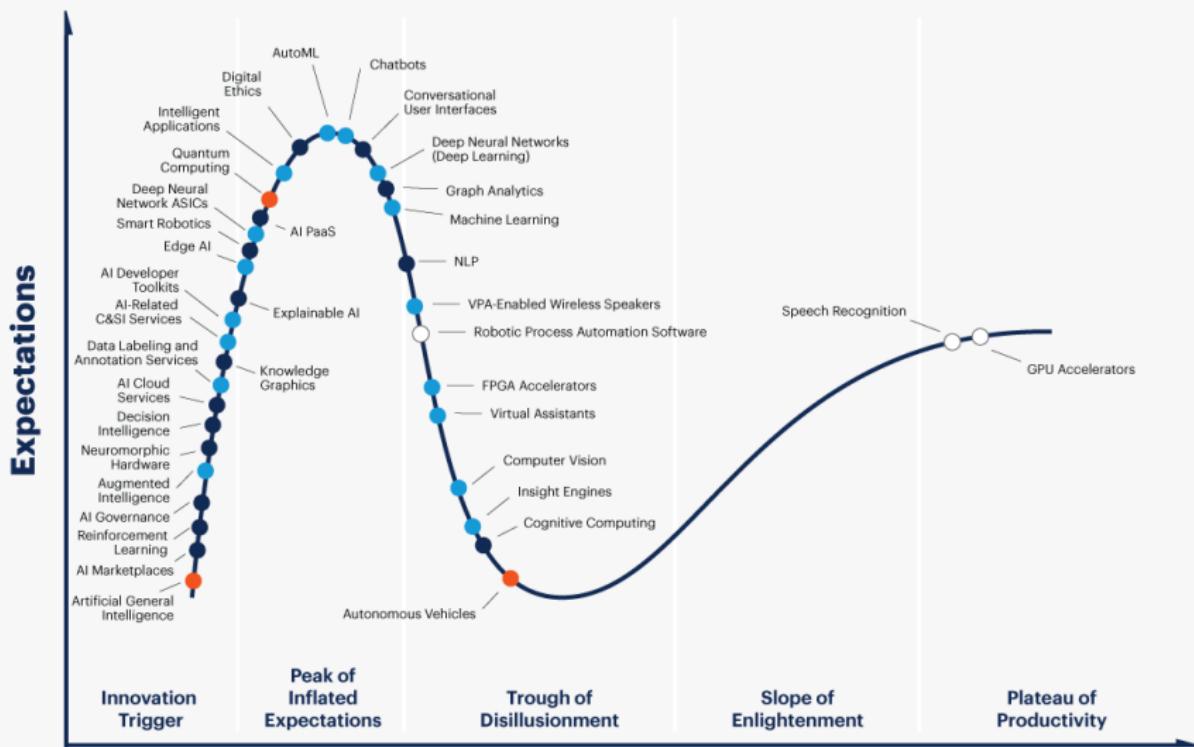


Gartner Hyper-Cycle of Emerging Technologies (2018)





Gartner Hyper-Cycle of Emerging Technologies (2019)





What is deep learning?

Deep learning has various closely related definitions or high-level descriptions.

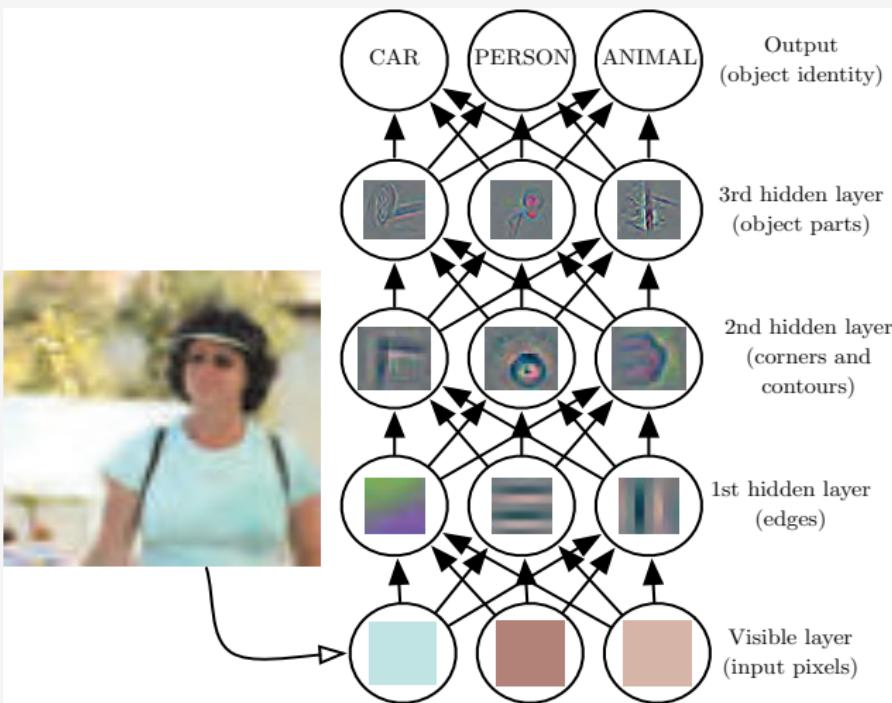
Definition (Deep learning)

A sub-field of machine learning that is based on

- learning several levels of representations, corresponding to a hierarchy of features or factors or concepts,
- where
 - higher-level concepts are defined from lower-level ones, and
 - the same lower-level concepts can help to define many higher-level concepts.



An Example





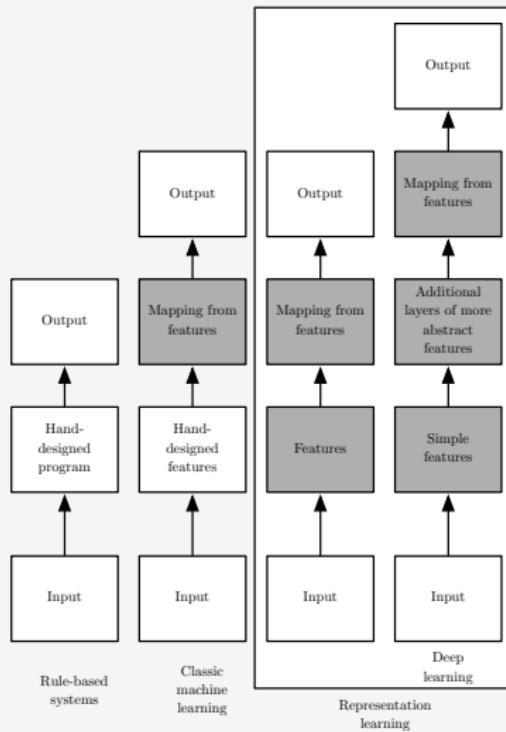
What is deep learning?

Definition (Deep learning)

- Deep learning is part of a broader family of machine learning methods based on **learning representations**.
- An observation (e.g., an image) can be represented in many ways (e.g., a vector of pixels), but some representations make it easier to learn tasks of interest (e.g., is this the image of a human face?) from examples, and research in this area attempts to define what makes better representations and how to learn them.



An Example





What is deep learning?

Common among the various high-level descriptions of deep learning are two key aspects:

- 1 Models consisting of **multiple layers/stages** of **nonlinear information processing**
- 2 Methods for supervised or unsupervised learning of feature representation at successively higher, more abstract layers.

Deep learning is in the intersections among the research areas of

- 1 Neural networks
- 2 Artificial intelligence
- 3 Graphical modeling
- 4 Optimization
- 5 Pattern recognition
- 6 Signal processing.



Table of contents

1 Course Information

2 Introduction

3 Success stories

4 Outline of course



Success stories¹

1 Finding nearest images



- day + night =



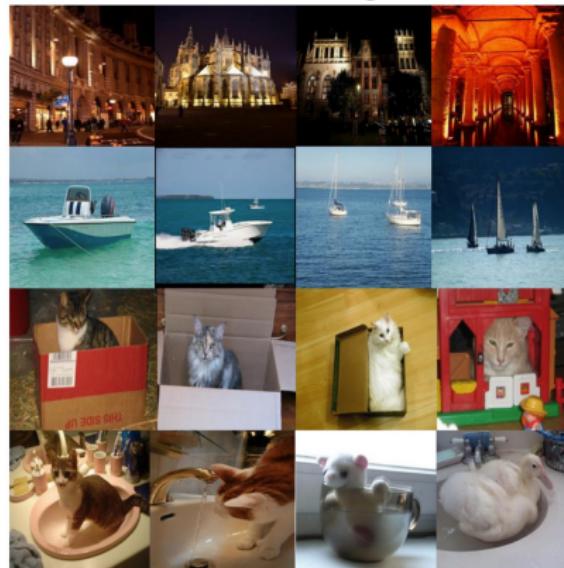
- flying + sailing =



- bowl + box =



- box + bowl =



¹This slide is taken from Prof. Ghodsi's slides.

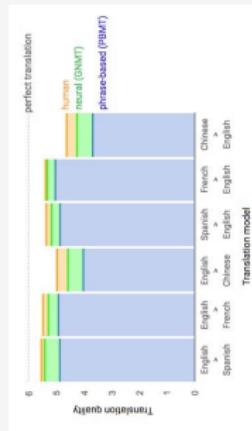


Success stories

- 1 Word2vec , Mikolov, 2013.

king man + woman = queen

- 2 Google neural machine translation²



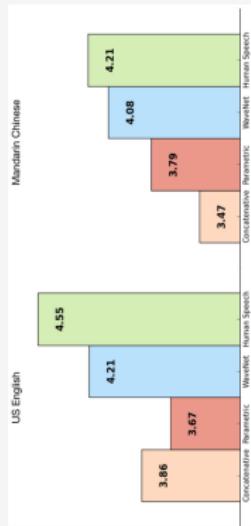
²Borrowed from

<https://blog.statmt.org/deep-learning-achievements-4c563e034257>



Success stories

1 Wavenet : Generating voice³



2 Lip Reading

³Borrowed from

<https://blog.statsbot.co/deep-learning-achievements-4c563e034257>



Success stories

1 LeNet-5

LeNet-5 is designed for handwritten and machine-printed character recognition

Live demo : <http://yann.lecun.com/exdb/lenet/index.html>

2 Sentiment Trees

Predicting the sentiment of movie reviews.

Live demo :

<http://nlp.stanford.edu:8080/sentiment/rntnDemo.html>



Success stories of Deep RL

- 1 TD-Gammon
- 2 DQN in Atari
- 3 Deep RL in Robotics
- 4 Alpha Go and Alpha Zero
- 5 Dota2 (Video Game)



Table of contents

1 Course Information

2 Introduction

3 Success stories

4 Outline of course



Outline of course

- 1 Introduction
- 2 Review of machine learning and history of deep learning
- 3 Multi-layer perceptrons and Backpropagation (MLP)
- 4 Optimization and Regularization
- 5 Convolutional networks (CNN)
- 6 Recurrent networks (RNN)
- 7 Sum-Product networks (SPN)
- 8 Dual learning
- 9 Deep reinforcement learning (Deep RL)
- 10 Representation learning
- 11 Deep generative models
- 12 Applications
 - Text mining and natural language processing
 - Computer vision
- 13 Advanced topics



Reading

Please read chapter 1 of Deep Learning Book.