

# Deep learning

## Word Embedding

Hamid Beigy

Sharif university of technology

November 22, 2019

# Table of contents

- 1 Introduction
- 2 Neural Probabilistic Language Model
- 3 Word2vec algorithm
- 4 Glove
- 5 Character-level embedding
- 6 Language model
- 7 Smart Reply
- 8 Reading

# Table of contents

- 1 Introduction
- 2 Neural Probabilistic Language Model
- 3 Word2vec algorithm
- 4 Glove
- 5 Character-level embedding
- 6 Language model
- 7 Smart Reply
- 8 Reading

# Introduction

## 1 How to represent a word?

- Represent words as atomic symbols such as talk, university, building.
- Represent word as a one-hot vector such as

$$\text{university} = \left( \underset{\text{egg}}{0}, \underset{\text{student}}{0}, \underset{\text{talk}}{0}, \underset{\text{university}}{1}, \underset{\text{building}}{0}, \dots, \underset{\text{buy}}{0} \right)$$

## 2 Issues with on-hot representation

- How large is this vector? dimensionality is large; vector is sparse
- Representing new words (any idea?).
- How measure word similarity?

# Distributional representation

- 1 Linguistic items with similar distributions have similar meanings  
(words occur in the same contexts probably have similar meaning).

$university = (0.2_{\text{egg}}, 0.1_{\text{student}}, 0.12_{\text{talk}}, 0.38_{\text{university}}, 0.2_{\text{building}}, \dots, 0.12_{\text{buy}})$

- 2 Word meanings are vector of basic concept.
- 3 What are basic concept?
- 4 How to assign weights?
- 5 How to define the similarity/distance?

# How to use word vectors?

## 1 Distance/similarity

- Cosine similarity: Word vector are normalized by length

$$\cos(u, v) = \frac{\langle u, v \rangle}{\|u\| \|v\|}$$

- Euclidean distance:

$$d(u, v) = \|u - v\|^2$$

- Inner product: This is same as cosine similarity if vectors are normalized

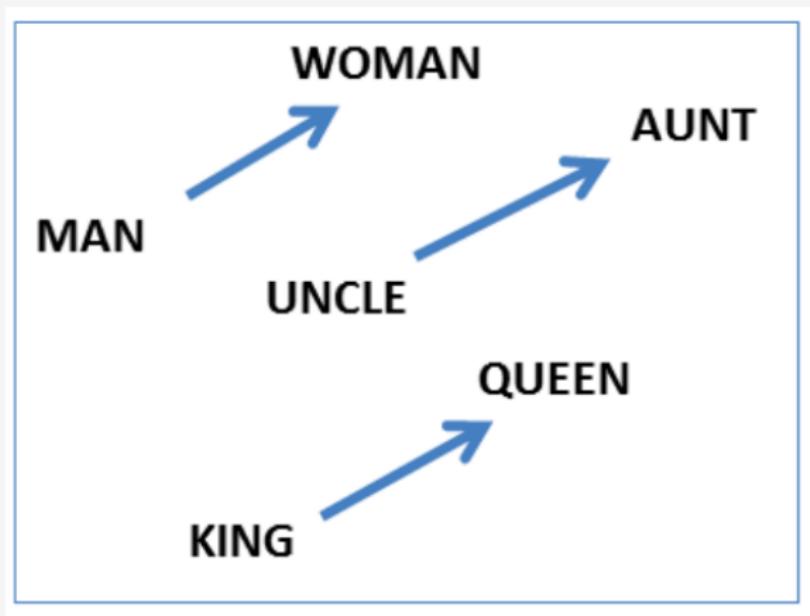
$$d(u, v) = \langle u, v \rangle$$

## 2 Choosing the right similarity metric is important.

# How to use word vectors?

## 1 Word analogy

$$v_{\text{man}} - v_{\text{woman}} + v_{\text{uncle}} \sim v_{\text{aunt}}$$



# How to learn word vectors?

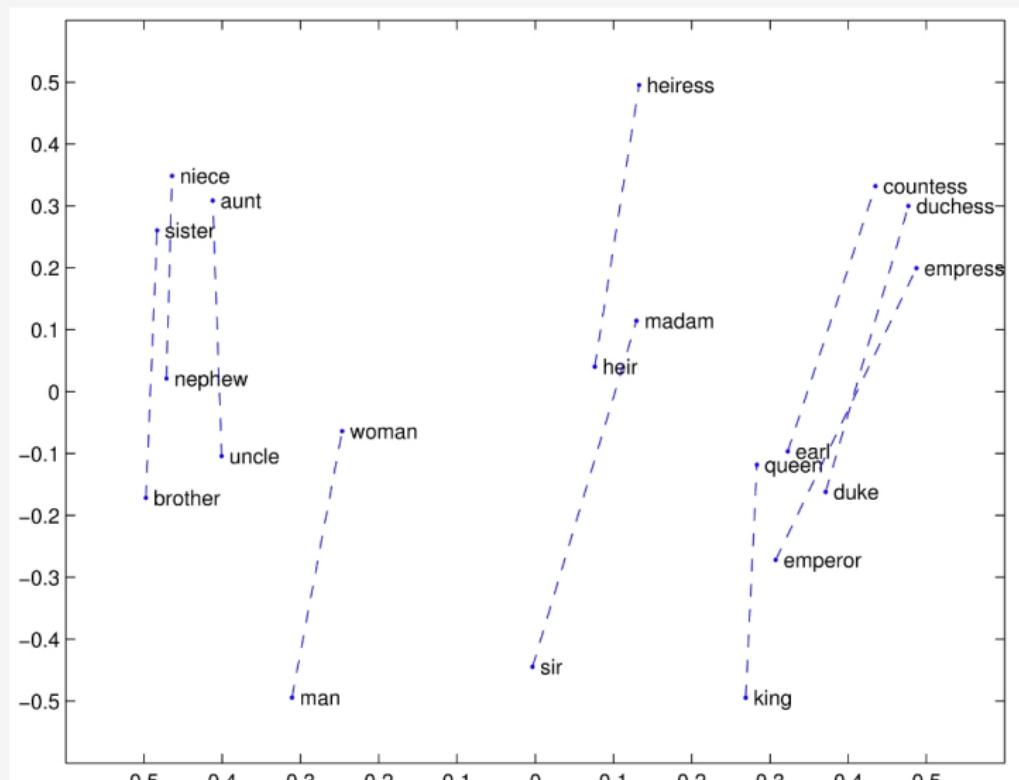
## 1 What are basic concept?

- We want that the number of basic concepts to be small and
- Basis be orthogonal

## 2 How to assign weights?

## 3 How to define the similarity/distance such as cosine similarity?

# Distributional representation (example)



# Term-document incidence matrix

## Example

	Anthony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth	...
Anthony	1	1	0	0	0	1	
Brutus	1	1	0	1	0	0	
Caesar	1	1	0	1	1	1	
Calpurnia	0	1	0	0	0	0	
Cleopatra	1	0	0	0	0	0	
mercy	1	0	1	1	1	1	
worser	1	0	1	1	1	0	
...							

Entry is 1 if term occurs. Example: Calpurnia occurs in *Julius Caesar*.

Entry is 0 if term doesn't occur. Example: Calpurnia doesn't occur in *Tempest*.

Each term is represented as a vector of bits.

# Term weighting

- 1 Evaluation of how important a term is with respect to a document.
- 2 First idea: the more important a term is, the more often it appears:  
*term frequency*

$$tf_{t,d} = \sum_{x \in d} f_t(x) \text{ where } f_t(x) = \begin{cases} 1 & \text{if } x = t \\ 0 & \text{otherwise} \end{cases}$$

- 3 The *order of terms* within a doc is ignored

# Inverse Document Frequency

- 1 *Inverse document frequency* of a term t:

$$idf_t = \log \frac{N}{df_t} \quad \text{with } N = \text{collection size}$$

- 2 Rare terms have high *idf*, contrary to frequent terms  
3 Example (Reuters collection):

Term t	df <sub>t</sub>	idf <sub>t</sub>
car	18165	1.65
auto	6723	2.08
insurance	19241	1.62
best	25235	1.5

- 4 In tf-idf weighting, the weight of a term is computed using both *tf* and *idf*:

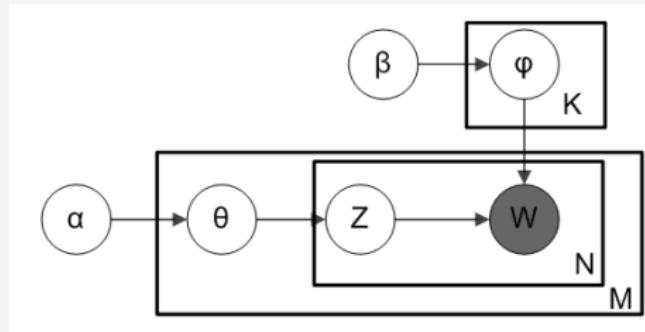
$$w(t, d) = tf_{t,d} \times idf_t \quad \text{called } tf - idf_{t,d}$$

# Dimensionality reduction

- 1 we don't need all of the dimensions that represent a word, only the most important ones.
- 2 There are several techniques such as
  - Principle Component Analysis (PCA): The most important dimensions contain the most variance
  - Latent Semantic Analysis (LSA): Project terms and documents into a topic space using SVD on term-document (co-occurrence) matrix.
  - Low-rank Approximation
- 3 Can we learn the dimensionality reduction from texts?

# Latent Dirichlet allocation

- 1 Assumes generative probabilistic model of a corpus<sup>1</sup>.
- 2 Documents are represented as distribution over latent topics, where each topic is characterized by a distribution over words.



---

<sup>1</sup>Blei, David M.; Ng, Andrew Y.; Jordan, Michael, "Latent Dirichlet Allocation". Journal of Machine Learning Research. 3 (45): pp. 993-1022, 2003.

# Table of contents

- 1 Introduction
- 2 Neural Probabilistic Language Model
- 3 Word2vec algorithm
- 4 Glove
- 5 Character-level embedding
- 6 Language model
- 7 Smart Reply
- 8 Reading

# Language modeling

- 1 An **language model** is a model for how humans **generate language**.
- 2 The quality of language models is measured based on their ability to learn a probability distribution over words in vocabulary  $V$ .
- 3 Language models generally try to compute the probability of a word  $w_t$  given its  $n - 1$  previous words, i.e.  $p(w_t|w_{t-1}, \dots, w_{t-n+1})$ .
- 4 Applying the chain rule and Markov assumption, we can approximate the probability of a whole sentence or document by the product of the probabilities of each word given its  $n$  previous words:

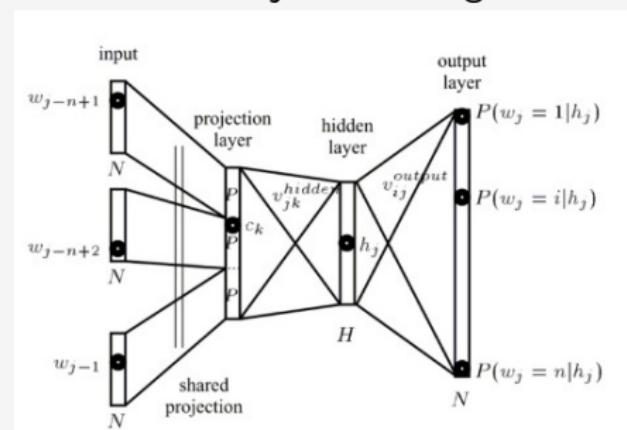
$$p(w_1, \dots, w_T) = \prod_i p(w_i|w_{i-1}, \dots, w_{i-n+1})$$

- 5 In n-gram based language models, we can calculate a word's probability based on the frequencies of its constituent n-grams:

$$p(w_t|w_{t-1}, \dots, w_{t-n+1}) = \frac{\text{count}(w_{t-n+1}, \dots, w_{t-1}, w_t)}{\text{count}(w_{t-n+1}, \dots, w_{t-1})}$$

# Neural Probabilistic Language Model

1 In NNs, we achieve the same objective using the softmax layer<sup>2</sup>.



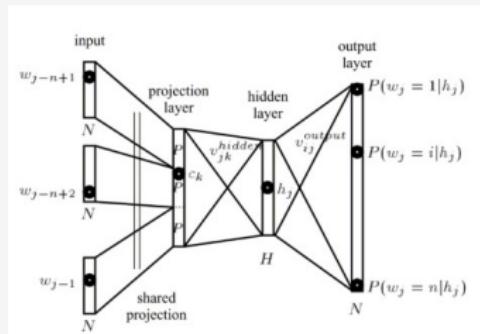
$$p(w_t | w_{t-1}, \dots, w_{t-n+1}) = \frac{\exp(h^\top v'_{w_t})}{\sum_{w_i \in V} \exp(h^\top v'_{w_i})}$$

---

<sup>2</sup>Bengio, Y., Ducharme, R., Vincent, P., and Janvin, C. (2003). A Neural Probabilistic Language Model. *The Journal of Machine Learning Research*, 3, 1137-1155.

# Neural Probabilistic Language Model

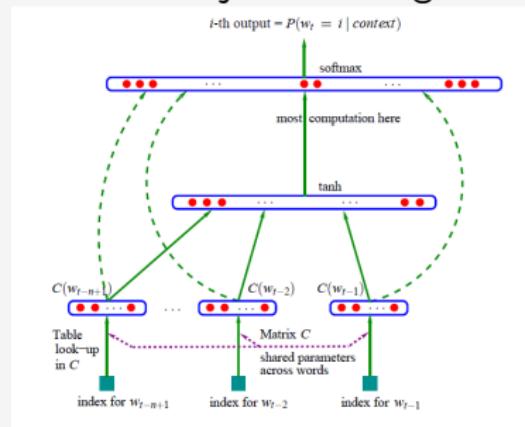
## 1 In this model



- 2 The inner product  $h^\top v'_{w_t}$  computes the (unnormalized) log-probability of word  $w_t$ , which we normalize by the sum of the log-probabilities of all words in  $V$ .
- 3  $h$  is the output vector of the penultimate network layer, while  $v'_{w_t}$  is the output embedding of word  $w_t$ , i.e. its representation in the weight matrix of the softmax layer.

# Neural Probabilistic Language Model

- 1 In NNs, we achieve the same objective using the softmax layer<sup>3</sup>.



- 2 Associate each word in vocabulary a distributed feature vector.
- 3 Learn both the embedding and parameters for probability function jointly.

<sup>3</sup>Bengio, Y., Ducharme, R., Vincent, P., and Janvin, C. (2003). A Neural Probabilistic Language Model. *The Journal of Machine Learning Research*, 3, 1137-1155.

# Table of contents

- 1 Introduction
- 2 Neural Probabilistic Language Model
- 3 Word2vec algorithm
- 4 Glove
- 5 Character-level embedding
- 6 Language model
- 7 Smart Reply
- 8 Reading

# Word2vec algorithm

- 1 Proposed by Mikolov et. al. and widely used for many NLP applications (in two papers).
- 2 Key features
  - Uses neural networks to train word / context classifiers (feedforward neural net)
  - Uses local context windows (environment around any word in a corpus) as inputs to the NN
  - Removed hidden layer.
  - Use of additional context for training LMs.
  - Introduced newer training strategies using huge database of words efficiently.

# Word2vec algorithm

- 1 In their first paper, Mikolov et al. propose two architectures for learning word embeddings that are computationally less expensive than previous models<sup>4</sup>.
- 2 In their second paper, they improve upon these models by employing additional strategies to enhance training speed and accuracy<sup>5</sup>.

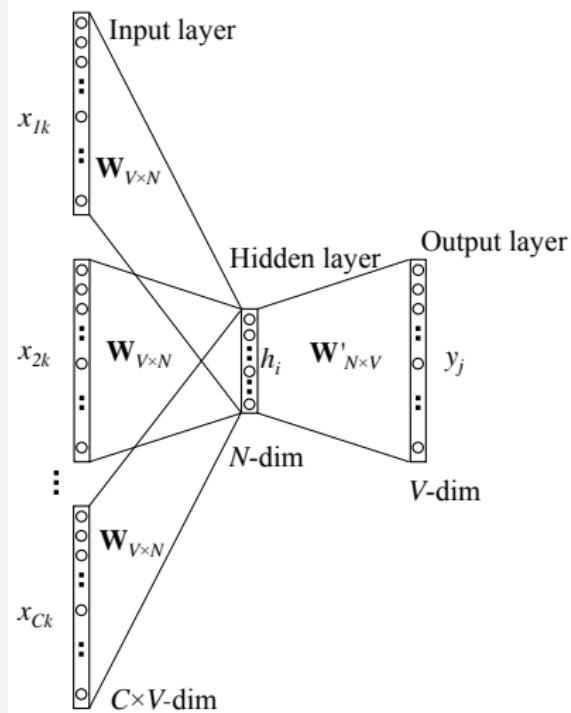
---

<sup>4</sup>Mikolov, T., Corrado, G., Chen, K., & Dean, J. (2013). Efficient Estimation of Word Representations in Vector Space. Proceedings of the International Conference on Learning Representations (ICLR 2013), 1-12.

<sup>5</sup>Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Distributed Representations of Words and Phrases and their Compositionality. NIPS, 1-9.

# Continuous Bag-of-Words

- 1 Mikolov et al. thus use both the  $n$  words before and after the target word  $w_t$  to predict it.
- 2 They call this continuous bag-of-words (CBOW), as it uses continuous representations whose order is of no importance.



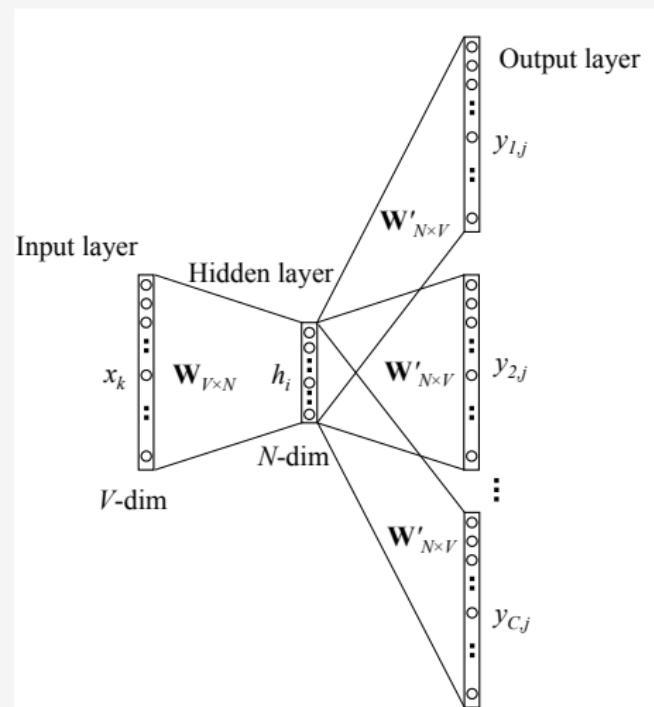
# Continuous Bag-of-Words objective function

- 1 The objective function of CBOW in turn is

$$I(\theta) = \sum_{t \in Text} \log P(w_t | w_{t-n}, \dots, w_{t-1}, w_{t+1}, \dots, w_{t+n})$$

# Skip-gram

- Instead of using the surrounding words to predict the center word as with CBOW, skip-gram uses the centre word to predict the surrounding words.



# Skip-gram objective function

- 1 The skip-gram objective thus sums the log probabilities of the surrounding  $n$  words to the left and to the right of the target word  $w_t$  to produce the following objective function.

$$I(\theta) = \sum_{t \in Text} \sum_{-n \leq j \leq n, j \neq 0} \log P(w_{t+j} | w_t)$$

# Table of contents

- 1 Introduction
- 2 Neural Probabilistic Language Model
- 3 Word2vec algorithm
- 4 Glove
- 5 Character-level embedding
- 6 Language model
- 7 Smart Reply
- 8 Reading

# Global Vectors for Word Representation<sup>6</sup>

- 1 Skip-gram doesn't utilize the statistics of corpus since they train on separate local context windows instead of on global co-occurrence counts.
- 2 The statistics of word occurrences in a corpus is the primary source of information available to all unsupervised methods for learning word representations.
- 3 Let  $X_{ij}$  be the number of times word  $j$  occurs in the context of word  $i$ .
- 4 Let  $X_i = \sum_k X_{ik}$  be the number of times any word appears in the context of word  $i$ .
- 5 Let  $P_{ij} = P(j|i) = \frac{X_{ij}}{X_i}$  be the probability that word  $j$  appear in the context of word  $i$ .

---

<sup>6</sup>Pennington, J., Socher, R., & Manning, C. D. (2014). *Glove: Global Vectors for Word Representation*. Proc. of Conference on Empirical Methods in Natural Language Processing, 1532-1543.

# Global Vectors for Word Representation

- 1 Co-occurrence probabilities for target words *ice* and *steam* with selected context words from a 6 billion token corpus.

Probability and Ratio	$k = \text{solid}$	$k = \text{gas}$	$k = \text{water}$	$k = \text{fashion}$
$P(k \text{ice})$	$1.9 \times 10^{-4}$	$6.6 \times 10^{-5}$	$3.0 \times 10^{-3}$	$1.7 \times 10^{-5}$
$P(k \text{steam})$	$2.2 \times 10^{-5}$	$7.8 \times 10^{-4}$	$2.2 \times 10^{-3}$	$1.8 \times 10^{-5}$
$P(k \text{ice})/P(k \text{steam})$	8.9	$8.5 \times 10^{-2}$	1.36	0.96

- 2 Considering two words  $i = \text{ice}$  and  $j = \text{steam}$  and study their relationship using various probe words,  $k$ .
- 3 For words  $k$  related to *ice* but *not steam*, say  $k = \text{solid}$ , we expect the ratio  $\frac{P_{ik}}{P_{jk}}$  will be large.
- 4 For words  $k$  related to *steam* but *not ice*, say  $k = \text{gas}$ , we expect the ratio  $\frac{P_{ik}}{P_{jk}}$  will be small.
- 5 For words  $k$  like *water* or *fashion*, that are either related to both *ice* and *steam*, or *to neither*, the ratio should be close to one.

# Global Vectors for Word Representation

- 1 This argument suggests that the appropriate starting point for word vector learning should be with ratios of co-occurrence probabilities rather than the probabilities themselves.
- 2 Noting that the ratio  $\frac{P_{ik}}{P_{jk}}$  depends on three words  $i, j, k$ , the most general model takes the form of

$$F(w_i, w_j, \tilde{w}_k) = \frac{P_{ik}}{P_{jk}}$$

where  $w \in \mathbb{R}^d$  are word vectors and  $\tilde{w} \in \mathbb{R}^d$  are separate context word vectors.

- 3 We would like  $F$  to encode  $\frac{P_{ik}}{P_{jk}}$  in the word vector space.
- 4 Since vector spaces are inherently linear structures, the most natural way to do this is with vector differences. Hence

$$F(w_i - w_j, \tilde{w}_k) = \frac{P_{ik}}{P_{jk}}$$

# Global Vectors for Word Representation

- 1 Parameters of  $F$  are vectors while the right-hand side is a scalar.
- 2  $F$  can be a complicated function such as a neural network, but a simplified function can be used also.

$$F((w_i - w_j)^T \tilde{w}_k) = \frac{P_{ik}}{P_{jk}}$$

- 3 For word-word co-occurrence matrices, the distinction between a word and a context word is arbitrary. Hence, we are free to exchange the two roles, i.e.  $w \leftrightarrow \tilde{w}$  and  $X \leftrightarrow X^T$ .
- 4 Hence, model should be invariant under this relabeling. Thus

$$F((w_i - w_j)^T \tilde{w}_k) = \frac{P_{ik}}{P_{jk}} = \frac{F(w_i^T \tilde{w}_k)}{F(w_j^T \tilde{w}_k)} = \frac{X_{ik}}{X_j}$$

- 5  $F = \exp$  is the solution of the above equation. Hence,

$$w_i^T \tilde{w}_k = \log P_{ik} = \log X_{ik} - \log X_i$$

# Global Vectors for Word Representation

- 1 We have,

$$\mathbf{w}_i^T \tilde{\mathbf{w}}_k = \log P_{ik} = \log X_{ik} - \log X_i$$

- 2 The above equation would exhibit the exchange symmetry if not for the  $\log X_i$  on the right-hand side.
- 3 This term is independent of  $k$  so it can be absorbed into a bias  $b_i$  for  $\mathbf{w}_i$ .
- 4 Adding an additional bias  $\tilde{b}_k$  for  $\tilde{\mathbf{w}}_k$  restores the symmetry. Hence

$$\mathbf{w}_i^T \tilde{\mathbf{w}}_k + b_i + \tilde{b}_k = \log X_{ik}$$

## Glove cost function

- 1 In an ideal setting, where you have perfect word vectors, the following expression will be zero.

$$w_i^T \tilde{w}_k + b_i + \tilde{b}_k - \log X_{ik} =$$

- 2 Hence, we can define the objective function as

$$J(w_i, \tilde{w}_k) = (w_i^T \tilde{w}_k + b_i + \tilde{b}_k - \log X_{ik})^2$$

- 3 Now, the final cost function is

$$J(w_i, \tilde{w}_k) = \sum_{i,k=1}^{|V|} f(X_{ik}) (w_i^T \tilde{w}_k + b_i + \tilde{b}_k - \log X_{ik})^2$$

where  $f$  is a weighting function, which is defined manually.

# Glove Results

1 GloVe becomes a global model for unsupervised learning of word representations that outperforms other models on word analogy, word similarity, and named entity recognition tasks.

## 2 Advantages

- Fast training
- Scalable to huge corpora
- Good performance even with small corpus, and small vectors
- Early stopping. We can stop training when improvements become small.

## 3 Drawbacks

- Uses a lot of memory: the fastest way to construct a term-cooccurrence matrix is to keep it in RAM as a hash map and perform cooccurrence increments in a global manner
- Sometimes quite sensitive to initial learning rate

# Table of contents

- 1 Introduction
- 2 Neural Probabilistic Language Model
- 3 Word2vec algorithm
- 4 Glove
- 5 Character-level embedding
- 6 Language model
- 7 Smart Reply
- 8 Reading

# FastText

- 1 One drawback of the Word2vec and Glove is the fact that they don't handle out-of-vocabulary.
- 2 FastText introduced the concept of subword-level embeddings, based on the skip-gram model, but where each word is represented as a bag of character -grams<sup>7</sup>.
- 3 A vector representation is associated to each character -gram, and words are represented as the sum of these representations.
- 4 This allows the model to compute word representations for words that did not appear in the training data.

---

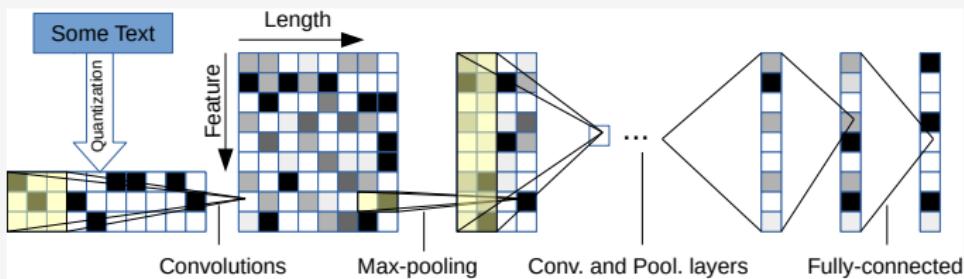
<sup>7</sup>Piotr Bojanowski, Edouard Grave, Armand Joulin, and Tomas Mikolov, [Enriching Word Vectors with Subword Information](#), Transactions of the Association for Computational Linguistics, Vol. 5, pp. 135–146, 2017.

# FastText

- 1 Each word  $w$  is represented as a bag of character  $n$ -gram, plus a special boundary symbols  $<$  and  $>$  at the beginning and end of words, plus the word  $w$  itself in the set of its  $n$ -grams.
- 2 Taking the word **where** and  $n = 3$  as an example, it will be represented by the character  $n$ -grams:  
**< wh, whe, her, ere, re> and the special sequence < where >.**
- 3 In practice, they extracted all the **n-grams** for  $n$  greater or equal to **3** and smaller or equal to **6**.

# Character-level Convolutional Networks<sup>8</sup>

- 1 A list of character are defined 70 characters which including 26 English letters, 10 digits, 33 special characters and new line character.
- 2 The network architectures are: 9 layers deep with 6 convolutional layers and 3 fully-connected layers.



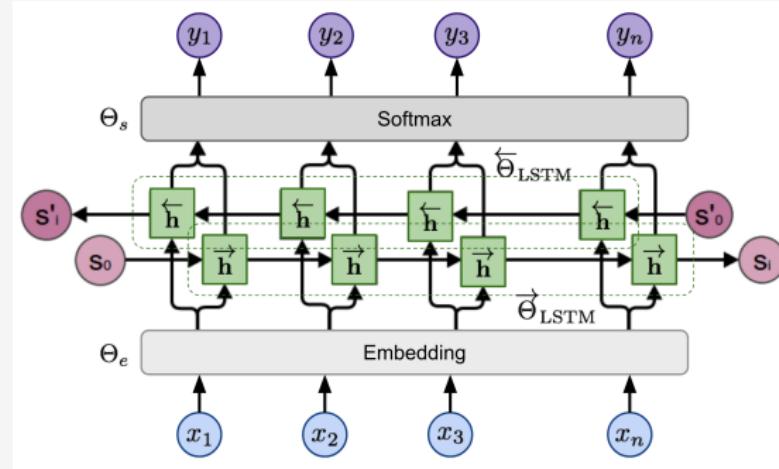
<sup>8</sup>Xiang Zhang, Junbo Zhao, and Yann LeCun, [Character-level Convolutional Networks for Text Classification](#), NIPS 2015.

# Table of contents

- 1** Introduction
- 2** Neural Probabilistic Language Model
- 3** Word2vec algorithm
- 4** Glove
- 5** Character-level embedding
- 6** Language model
- 7** Smart Reply
- 8** Reading

# Embeddings from Language Model (ELMo) <sup>9</sup>

- 1 ELMo learns contextualized word representation by pre-training a language model in an unsupervised way.



<sup>9</sup>Matthew Peters, et. al., Deep Contextualized Word Representations, Proc. of Conf. of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers), 2018.

# Embeddings from Language Model (ELMo)

- 1 The bidirectional Language Model (biLM) is the foundation for ELMo.
- 2 While the input is a sequence of  $n$  tokens,  $(x_1, \dots, x_n)$ , the language model learns to predict the probability of next token given the history.
- 3 In the forward pass, the history contains words before the target token,

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i \mid x_1, \dots, x_{i-1})$$

- 4 In the backward pass, the history contains words after the target token,

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i \mid x_{i+1}, \dots, x_n)$$

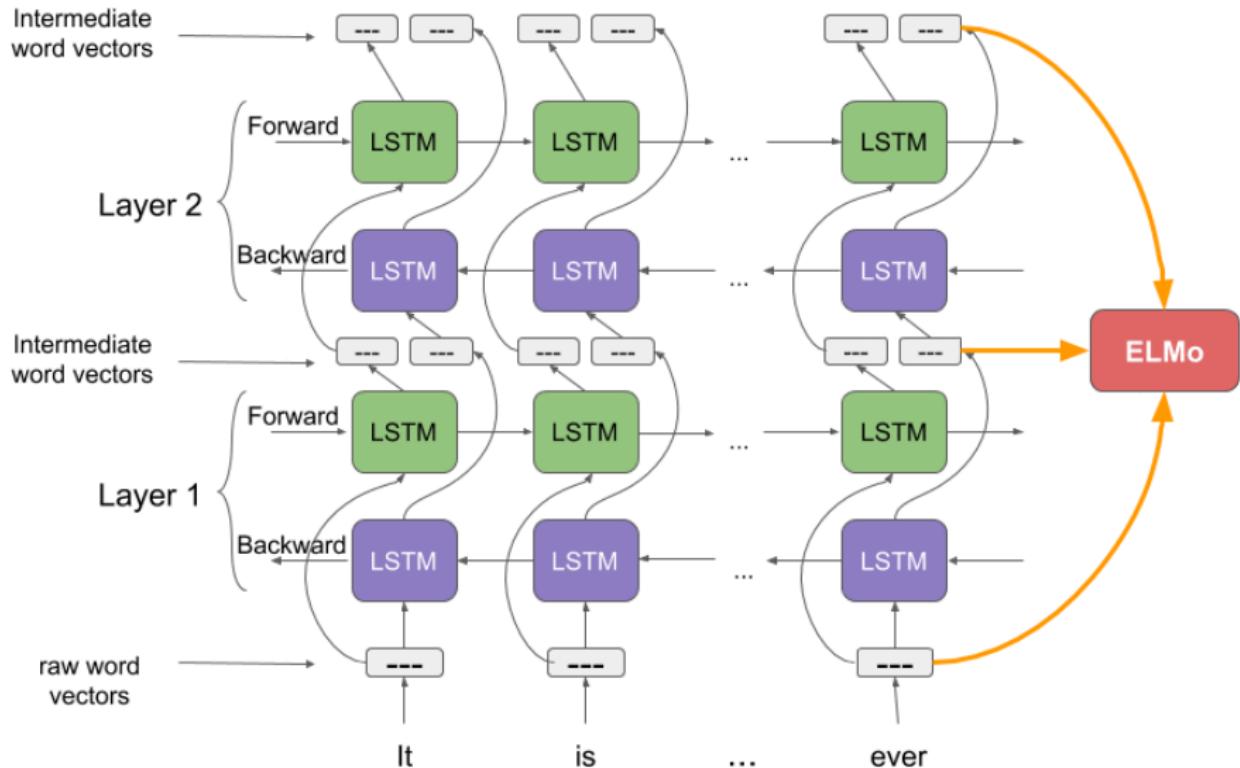
- 5 The predictions in both directions are modeled by multi-layer LSTMs with hidden states.

# Embeddings from Language Model (ELMo)

- 1 The model is trained to minimize the negative log likelihood (= maximize the log likelihood for true words) in both directions:

$$\mathcal{L} = - \sum_{i=1}^n \left( \log p(x_i \mid x_1, \dots, x_{i-1}; \Theta_e, \vec{\Theta}_{\text{LSTM}}, \Theta_s) + \log p(x_i \mid x_{i+1}, \dots, x_n; \Theta_e, \overleftarrow{\Theta}_{\text{LSTM}}, \Theta_s) \right)$$

# Embeddings from Language Model (ELMo)

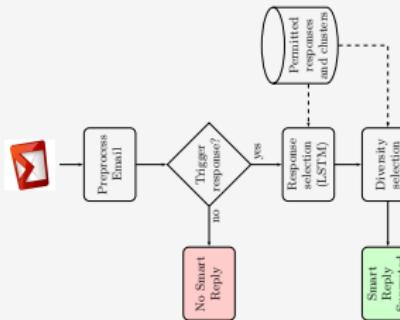


# Table of contents

- 1** Introduction
- 2** Neural Probabilistic Language Model
- 3** Word2vec algorithm
- 4** Glove
- 5** Character-level embedding
- 6** Language model
- 7** Smart Reply
- 8** Reading

# Smart Reply<sup>10</sup>

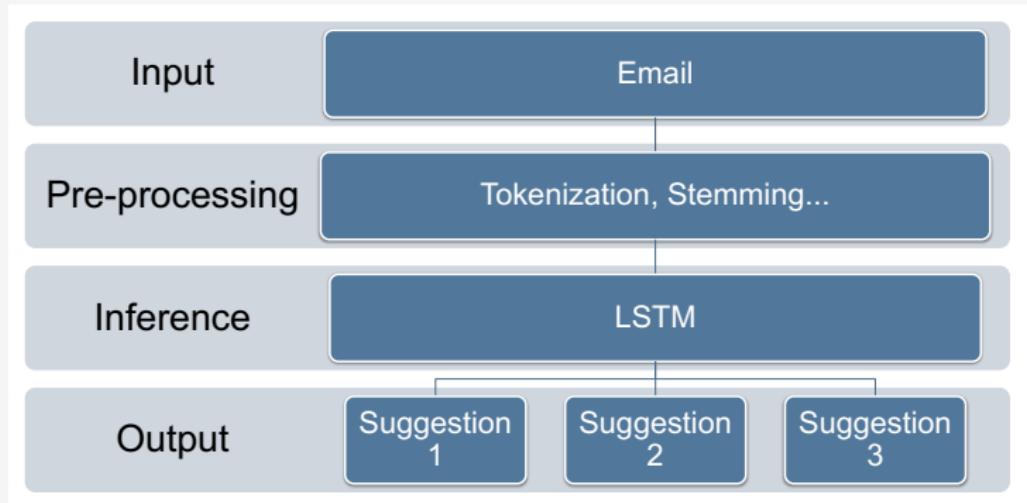
- 1 An end-to-end method for automatically generating short email responses, called **Smart Reply**.
- 2 It generates semantically diverse suggestions that can be used as complete email responses with just one tap on mobile.
- 3 The system is currently used in In- box by Gmail and is responsible for assisting with 10% of all mobile responses.



<sup>10</sup>Anjuli Kannan et. al., **Smart Reply: Automated Response Suggestion for Email**, SIGKDD 2016.

# Smart Reply Architecture

## 1 A first version



# Smart Reply Architecture

## 1 Issues Response diversity, Quality control, and Computing costs

### Unnormalized Responses

Yes, I'll be there.  
Yes, I will be there.  
I'll be there.  
Yes, I can.  
What time?  
I'll be there!  
I will be there.  
Sure, I'll be there.  
Yes, I can be there.  
Yes!

# Smart Reply Architecture

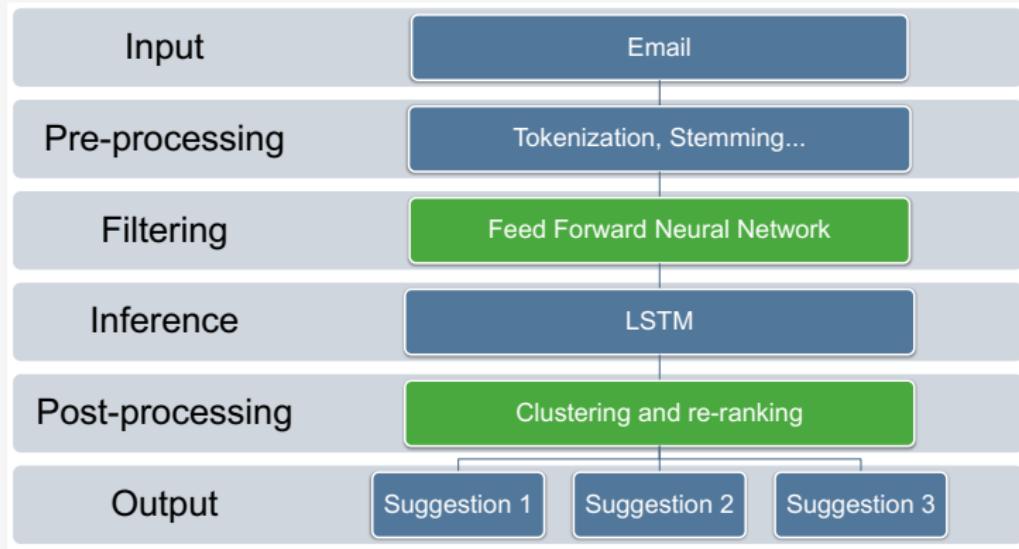
## 1 Issues Response diversity, Quality control, and Computing costs

Unnormalized Responses
Yes, I'll be there.
Yes, I will be there.
I'll be there.
Yes, I can.
What time?
I'll be there!
I will be there.
Sure, I'll be there.
Yes, I can be there.
Yes!

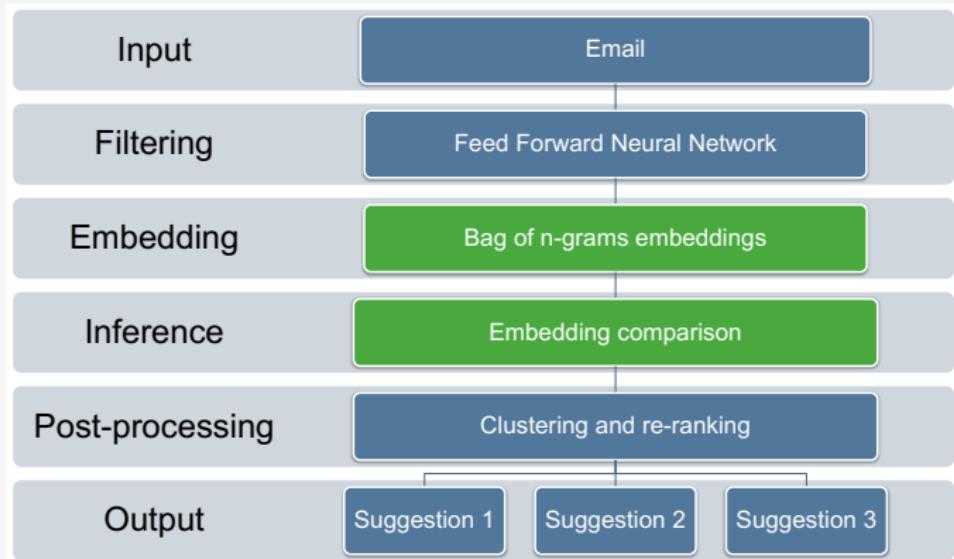
# Smart Reply: A better architecture

## 1 Issues

- Response diversity : semi supervised clustering (human in the loop)
- Quality control : restrict output sentences to a subset
- Computing costs : triggering NN system eliminates 90% of messages



# Smart Reply: An optimized version<sup>11</sup>



<sup>11</sup> Matthew Henderson, et al., [Efficient Natural Language Response Suggestion for Smart Reply](#), arXiv 2017.

# Table of contents

- 1 Introduction
- 2 Neural Probabilistic Language Model
- 3 Word2vec algorithm
- 4 Glove
- 5 Character-level embedding
- 6 Language model
- 7 Smart Reply
- 8 Reading

# Reading I

- 1 Blei, David M.; Ng, Andrew Y.; Jordan, Michael, "Latent Dirichlet Allocation". *Journal of Machine Learning Research*. 3 (45): pp. 993-1022, 2003.
- 2 Bengio, Y., Ducharme, R., Vincent, P., and Janvin, C. (2003). A Neural Probabilistic Language Model. *The Journal of Machine Learning Research*, 3, 1137-1155.
- 3 Mikolov, T., Corrado, G., Chen, K., & Dean, J. (2013). Efficient Estimation of Word Representations in Vector Space. Proc. of the International Conference on Learning Representations (ICLR 2013), 1-12.
- 4 Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Distributed Representations of Words and Phrases and their Compositionality. NIPS, 1-9.

# Reading II

- 5 Pennington, J., Socher, R., & Manning, C. D. (2014). Glove: Global Vectors for Word Representation. Proc. of the 2014 Conference on Empirical Methods in Natural Language Processing, 1532-1543.
- 6 Di Carlo, V., Bianchi, F., and Palmonari, M. (2019). Training Temporal Word Embeddings with a Compass. Proc. of the AAAI Conference on Artificial Intelligence, 33(01), 6326-6334.
- 7 White board notes.