

بِسْمِ اللّٰهِ الرَّحْمٰنِ الرَّحِيْمِ

Algorithmic Game Theory

Javad Salimi

Salimi.sartakhti@gmail.com

۱۳۹۶



Let's solve prisoners' dilemma!

REPEATED GAMES

Introduction

- We focus on a class of games with several interactions, i.e., **Repeated Game**
- Our ultimate goal is to model and solve the problem of prisoners' dilemma with a repeated interaction, examples include
 - Several *real life situations* such as, friendship, marriage, and wars.
 - *Engineering Applications* such as, multiple access protocols in wireless communications, packet forwarding, and jamming.

A Brief Reminder

- A Nash Equilibrium $(s_1^*, s_2^*, \dots, s_N^*)$ is a **Sub-Game Perfect Nash Equilibrium (SPNE)** if it induces a Nash Equilibrium in every sub-game of the game
- We looked for the Nash equilibria in each of the sub-games, roll the payoffs back up, and then see what the optimal moves are higher up the tree (e.g., Mixed NE in the BoS Game)

The War of Attrition

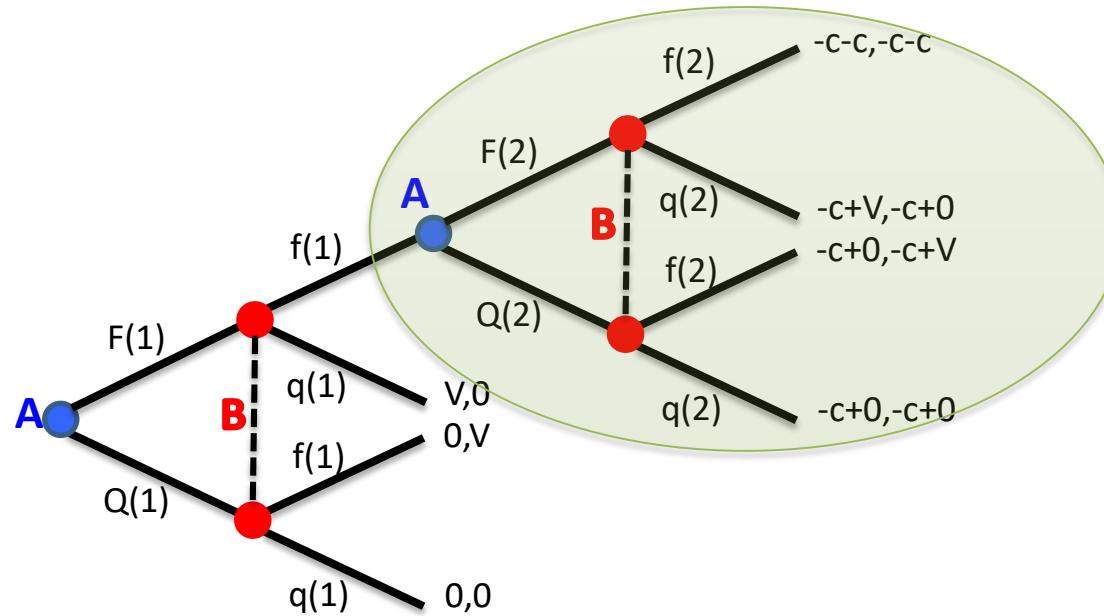
- 2-Player game
- In each period of the game each chooses **Fight (F)** or **Quit (Q)**
 1. If the other player quits first, you win a prize **V**
 2. Each period in which both **F**, each player pays cost **-c**
 3. If both quit at once they get **0**

(Repeated TDMA Transmission)

- Example: First World War, The Battle for Broadband [IEEE Spectrum, 2005]

Let's Play This Game Repeatedly! 5

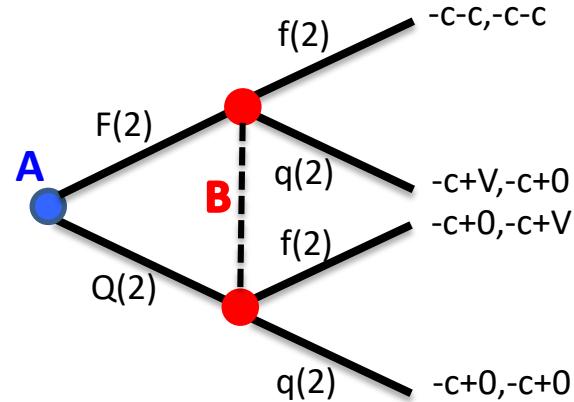
Two Period Game



Assumptions:

- We now focus on a two stage game
- Later we will play it for infinite stages
- We assume that $v > c$

Analysis of the Second Sub-game

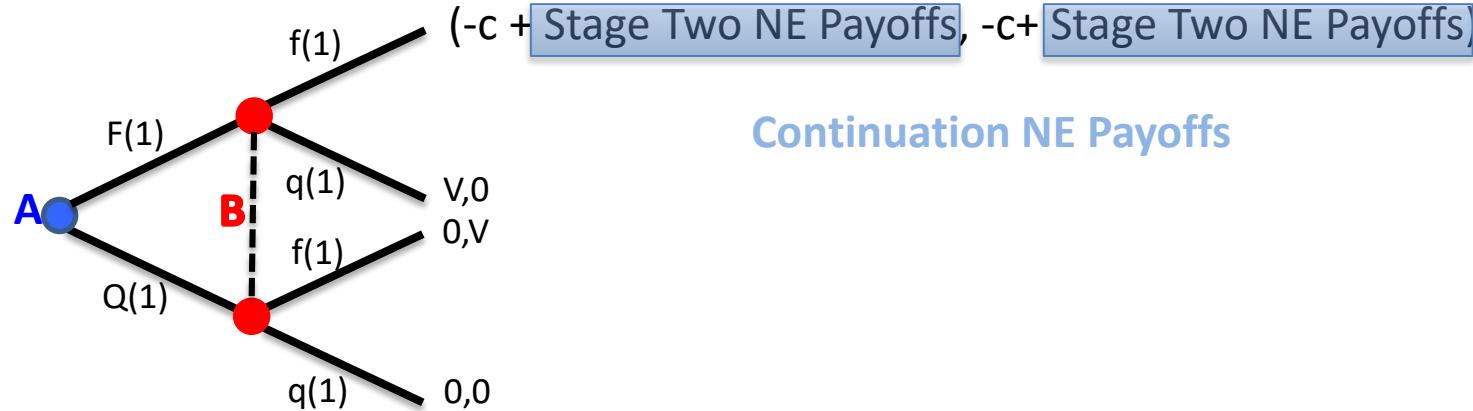


-C + A

		B	
		f(2)	q(2)
		F(2)	V, 0
		-C, -C	
		0, V	0, 0

- Two Pure-Strategy NE in this sub-game
 - $(F(2), q(2))$ and $(Q(2), f(2))$
 - Payoffs are $(V, 0)$ and $(0, V)$
- Note that c (sunk costs) does not matter!

Roll the Payoffs Back Up to the First Sub-Game



		B	
		f(1)	q(1)
A	F(1)	-c + V, -c + 0	V, 0
	Q(1)	0, V	0, 0

For $(F(2), q(2))$ in stage 2 (tomorrow)
The NE is: $(F(1), q(1))$

		B	
		f(1)	q(1)
A	F(1)	-c + 0, -c + V	V, 0
	Q(1)	0, V	0, 0

For $(Q(2), f(2))$ in stage 2 (tomorrow)
The NE is $(Q(1), f(1))$

Pure Strategy SPNE ($V > c$)

- Two SPNE (Fighter vs Quitter)
 - $[(F(1), F(2)), (q(1), q(2))]$
 - $[(Q(1), Q(2)), (f(1), f(2))]$
- **Main Lesson:** If we know I am going to win tomorrow, then I win today
- **Our first intuition:** Rational players should involve fights in NE, but here one only involves to fight and other quits!
- What did we miss?

Second Sub-game Analysis

Mixed Strategy NE

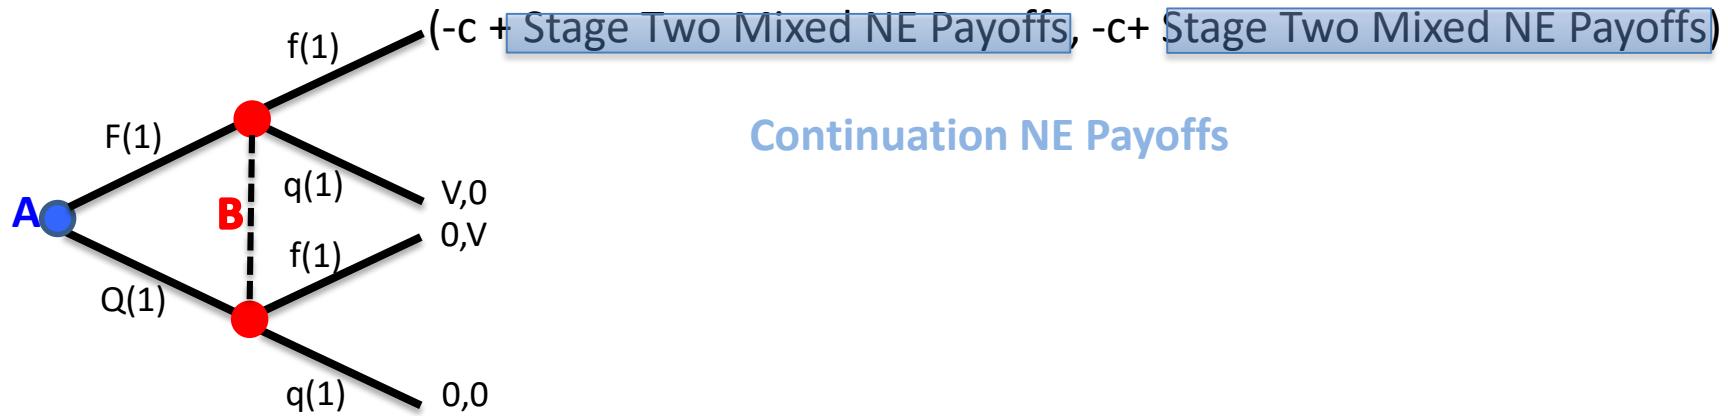
		B	
		f(2)	q(2)
		F(2)	-C, -C V, 0
		Q(2)	0, V 0, 0
		p	1-p

- If A Fights $\rightarrow -cp + V(1-p)$
- If A Quits $\rightarrow 0p + 0(1-p)$

$$V(1-p) = pc \rightarrow p = V/(V+c) \text{ and } 1-p = c/(V+c)$$

- Mixed NE has both fight with probability $= V/(V+c)$
- Payoffs in this mixed NE = (0,0)
- Probability of fight increases in V and decreases in c

Roll the Payoffs Back Up to the First Sub-Game



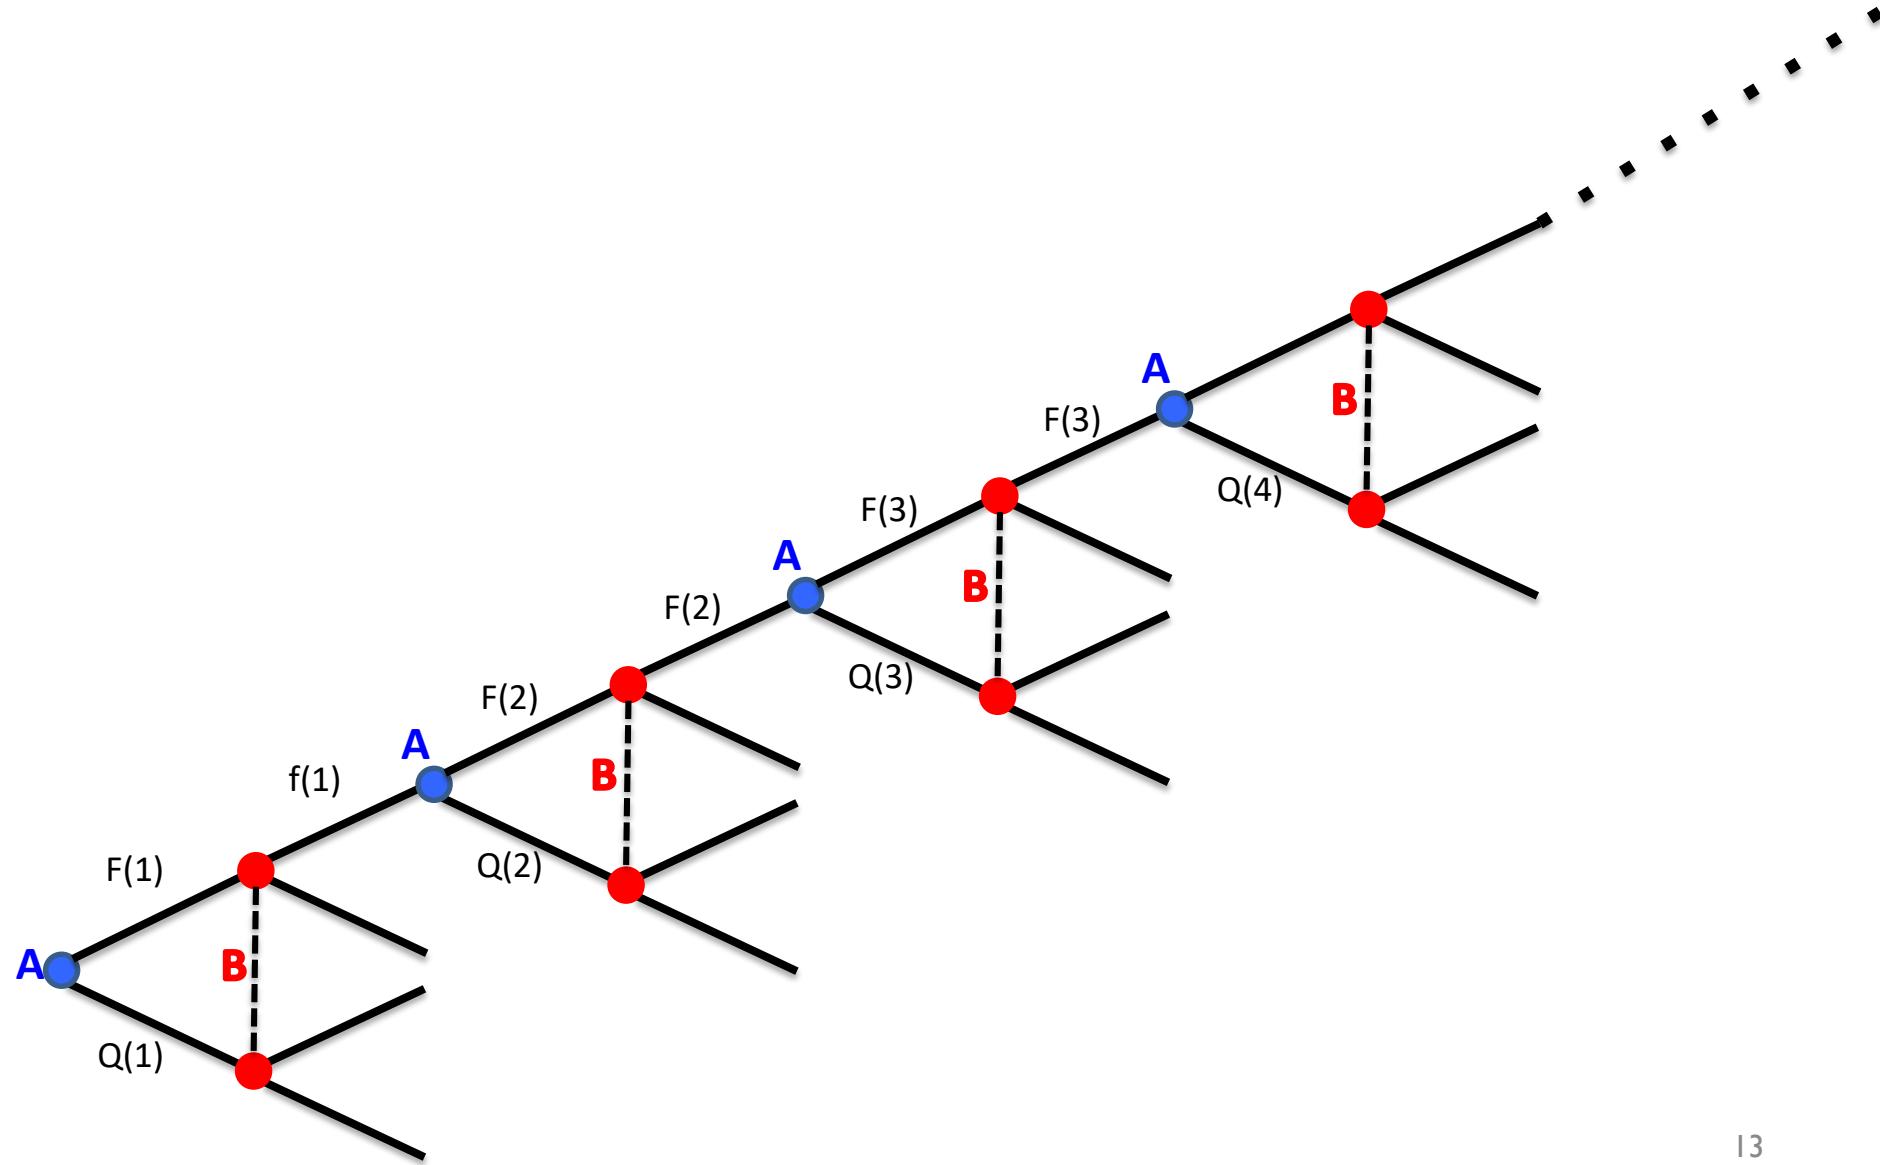
		$f(1)$	B	$q(1)$
A		$F(1)$	$-c + 0, -c + 0$	$V, 0$
A	$F(1)$	$0, V$		$0, 0$
	$Q(1)$			

← For the mixed NE in Period 2

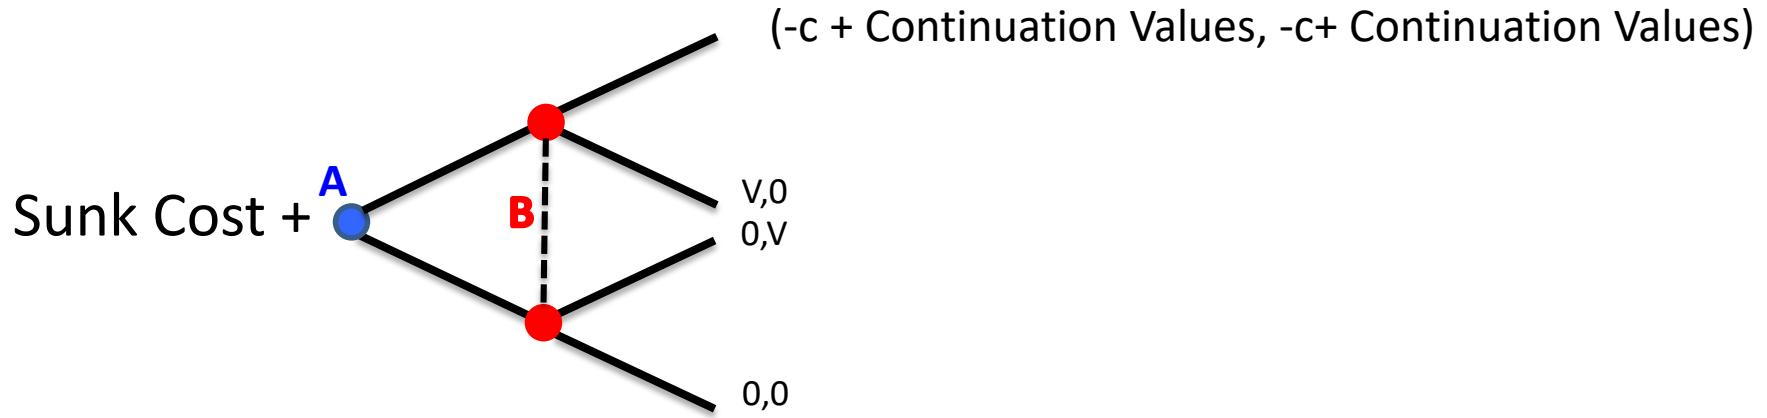
Conclusions

- When we rolled back, the matrix is the same as the second game (last sub-game)
- Same payoff matrix, so ...
- Both Fights with $p=V/(V+c)$ \rightarrow Mixed SPE $[(p^*, p^*), (p^*, p^*)]$
 \rightarrow Expected payoff is 0
- We end up fighting in each game with probability of p^* (in infinite horizon)

Let's Play Forever



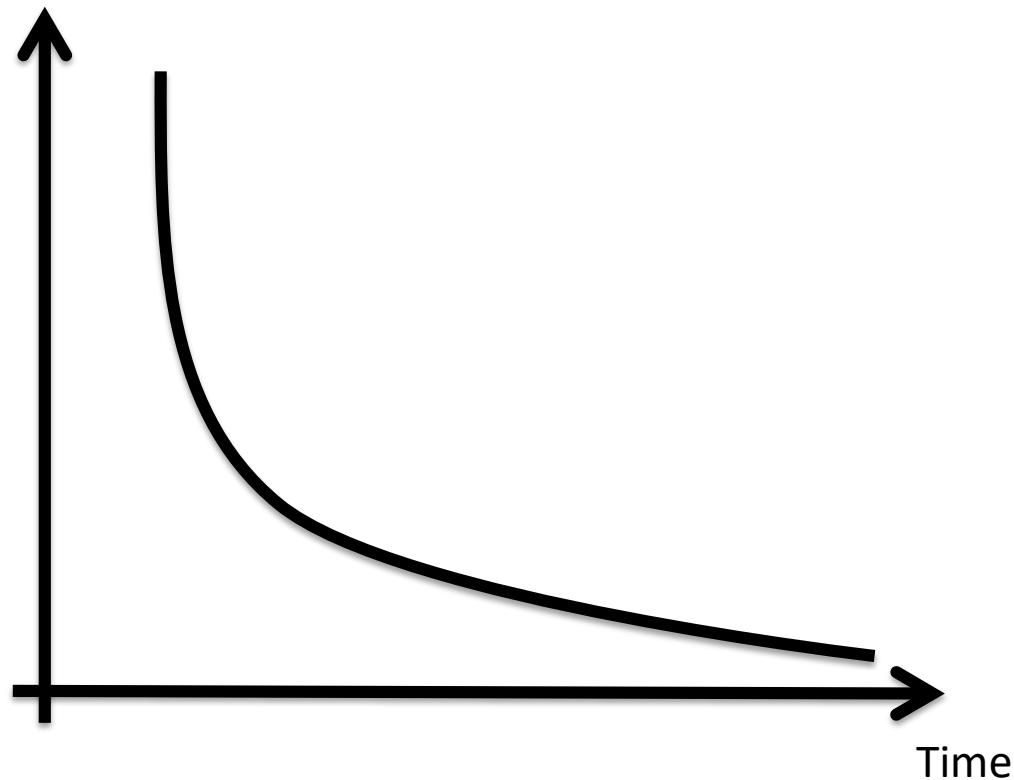
And Solution!



- ✧ Assume that we play in Stage “n” = 659870
- ✧ If mix in future then the continuation value is $(0,0)$
- ✧ This game is already analyzed (solved for mixed)
- ✧ Both mix with F with $p^* = V/(V+c)$
- ✧ And the fight will continue ...

Probability of Continued Fight

Probability of Continued Fight (War)



Long fight between rational players

Take away Messages

- We could sustain fighting by players that were **rational in a war of attrition.**
 - Break the analysis up into what we might call "stage games"
 - Break the payoffs up into: the payoffs that are associated with that stage and
 - Payoffs that are associated with the past (they're sunk, so they don't really matter)
 - Payoffs that are going to come in the future from future equilibrium play

REPEATED PRISONERS' DILEMMA

Relationships are Repeated Not Contractual

- **Friendship**
 - If you are nice to me I will nice to you
- **Nations Relationships**
 - Visa issues
- **Exchange goods and Services**
 - Change fruit with petrol

Why Repeated Interactions?

In ongoing relationships the promise of **future rewards** and the **threat of future punishments** may sometimes provide **incentives for good behavior today**

Good News: Repeated interaction might get us out of prisoners' dilemma

Prisoners' Dilemma

	B	
	Coop	Defect
A	Coop	2,2
	Defect	-1,3
	3,-1	0,0

Let's Play This Game Repeatedly!

Repeated Solution

- The game at the last stage is just a simple Prisoners' Dilemma
- Then player should defect
- We put payoffs of tomorrow in the game of today and we have

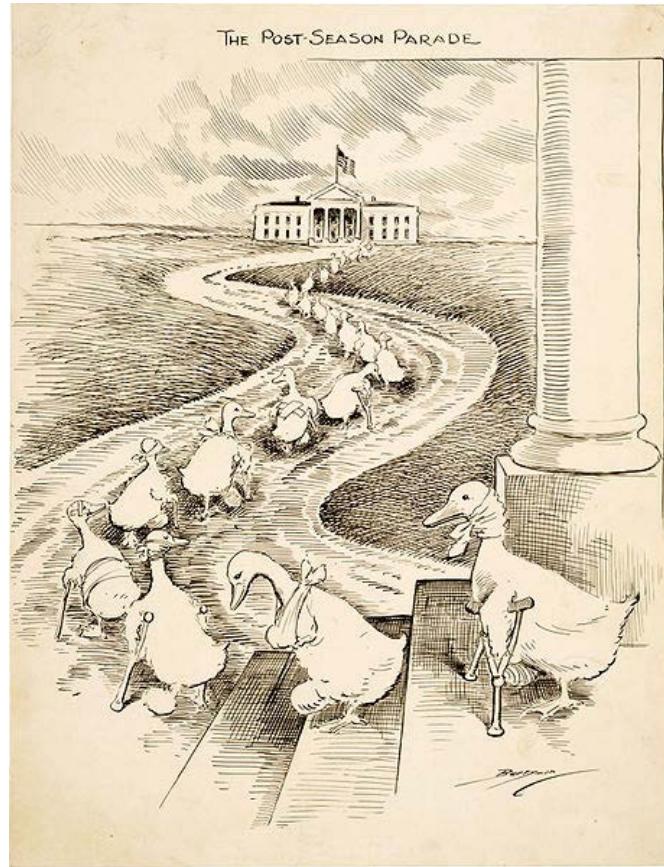
		B	
		Coop	Defect
		Coop	2+0,2+0
A	Coop	-1+0,3+0	
	Defect	3+0,-1+0	0+0,0+0

Repeated Solution

- It is again a PD game and we have again unraveling from back
- The problem is not solved!!!!

Example:

- Lame Duck Effect



- However, even a Finite game has some hope!
Let's discuss an example

Another Example

	A	B	C
A	4,4	0,5	0,0
B	5,0	1,1	0,0
C	0,0	0,0	3,3

- We play twice
- We would like to sustain (A,A) Cooperation
- But (A,A) is not a NE in one-shot game
- The NE are (B,B) and (C,C)
 - There are some Mixed NE, but let's focus on pure NE
- ➔ We cannot sustain (A,A) in period 2

New Strategy

- **Play A, then:**
 - Play C if (A,A) was played
 - Play B otherwise
- Pay attention to *information sets*
 - It says what to do in all 1+9 information sets (we are fine!)
- **Question: Is this a SPNE?**
- Or can we sustain Nash behavior in all subgames?

Repeated Analysis

- For each activity in the first period we make a new sub-game
- In period 2:
 - After (A,A) this strategy induces (C,C)
 - Which is a NE \circlearrowright
 - After the other choices in period 1, this strategy induces (B,B)
 - Which is a NE \circlearrowright

We are fine with all subgames after the first move
(i.e., 9 subgames)

What about the Whole Game?

- In the whole game:
 - A → $u(A,A) + u(C,C) = 4 + 3 = 7$
 - If Defect (an obvious defection, check the rest at home😊):
 - B → $u(B,A) + u(B,B) = 5 + 1 = 6$

*Temptation to defect (cheat) today ≤
(Value of rewards tomorrow – Value of punishment tomorrow)*

$$5 - 4 \leq 3 - 1 \rightarrow 1 \leq 2$$

→ Temptation is outweighed by the difference between the value of the reward and the value of the punishment

Important Lesson

If a “Stage Game” has more than one NE, then we **may be able** to use the prospect of playing different equilibria tomorrow to **provide incentives** (rewards and punishments) **for cooperating today**

A Brief Comment!

- There may be a problem of renegotiation
 - Between two stages they negotiate to switch to (C,C) → Then there is no incentive to cooperate in the first stage!
 - E.g., 2008 crisis (bail out or bankruptcy)

Repeated Prisoners' Dilemma

	B	
	Coop	Defect
A	Coop	2,2
	Defect	-1,3
	3,-1	0,0

- Each round we toss a coin twice and decide to follow the game or not
- Main difference: We do not know the end

Grim Trigger Strategy

		B	
		Coop	Defect
A	Coop	2,2	-1,3
	Defect	3,-1	0,0

- The game continue with probability of δ (*toss a coin two times*)
- **Play C then**
 - Play C if no one has played D
 - Play D otherwise

Is it NE?

- *Temptation to defect (cheat) today ≤ (Value of rewards tomorrow – Value of punishment tomorrow)*
 - $3-2 \leq \delta (u(C,C) \text{ forever} - u(D,D) \text{ forever})$
- Why δ ?
- Because the game may end
 - You need money today!

Is it NE?

$$3 - 2 \stackrel{?}{\leq} \delta (u(C,C)_{\text{forever}} - u(D,D)_{\text{forever}})$$

- $u(D,D)_{\text{forever}} = 0$
- $u(C,C)_{\text{forever}} = 2 + 2\delta + 2\delta^2 + 2\delta^3 + \dots$
 $= 2/(1-\delta)$

Is it NE?

- Is Grim Trigger a Nash Equilibrium?

$$I \leq [2/(I-\delta)-0] \delta$$

$$\delta \geq I/3$$

Any other NE?

- What about playing D now, then C, then D forever?
 - (D,C), (C,D), (D,D), (D,D), ...
 - $u = 3 + (-1) \delta + 0 + 0 + 0 + \dots = 3 - \delta$
 - It is even worse comparing to all D after the first D (D,D,D,...)
- Punishment (D,D) forever is a SPNE

Lesson

We can get cooperation in Prisoner's Dilemma using the *Grim Trigger Strategy* (as a SPNE) provided $\delta \geq 1/3$

General Lesson

For an ongoing relationship to provide incentives for good behavior today, it helps for there to be a high probability that the relationship will continue
(weight you put on the future)

One Period Punishment

- Play C to start, then
 - Play C if either (C,C) or (D,D) were played last
 - Play D if either (C,D) or (D,C) were played last

Is this a SPNE?

- *Temptation to defect (cheat) today* \leq
*(Value of promise tomorrow – Value of the threat
tomorrow)*

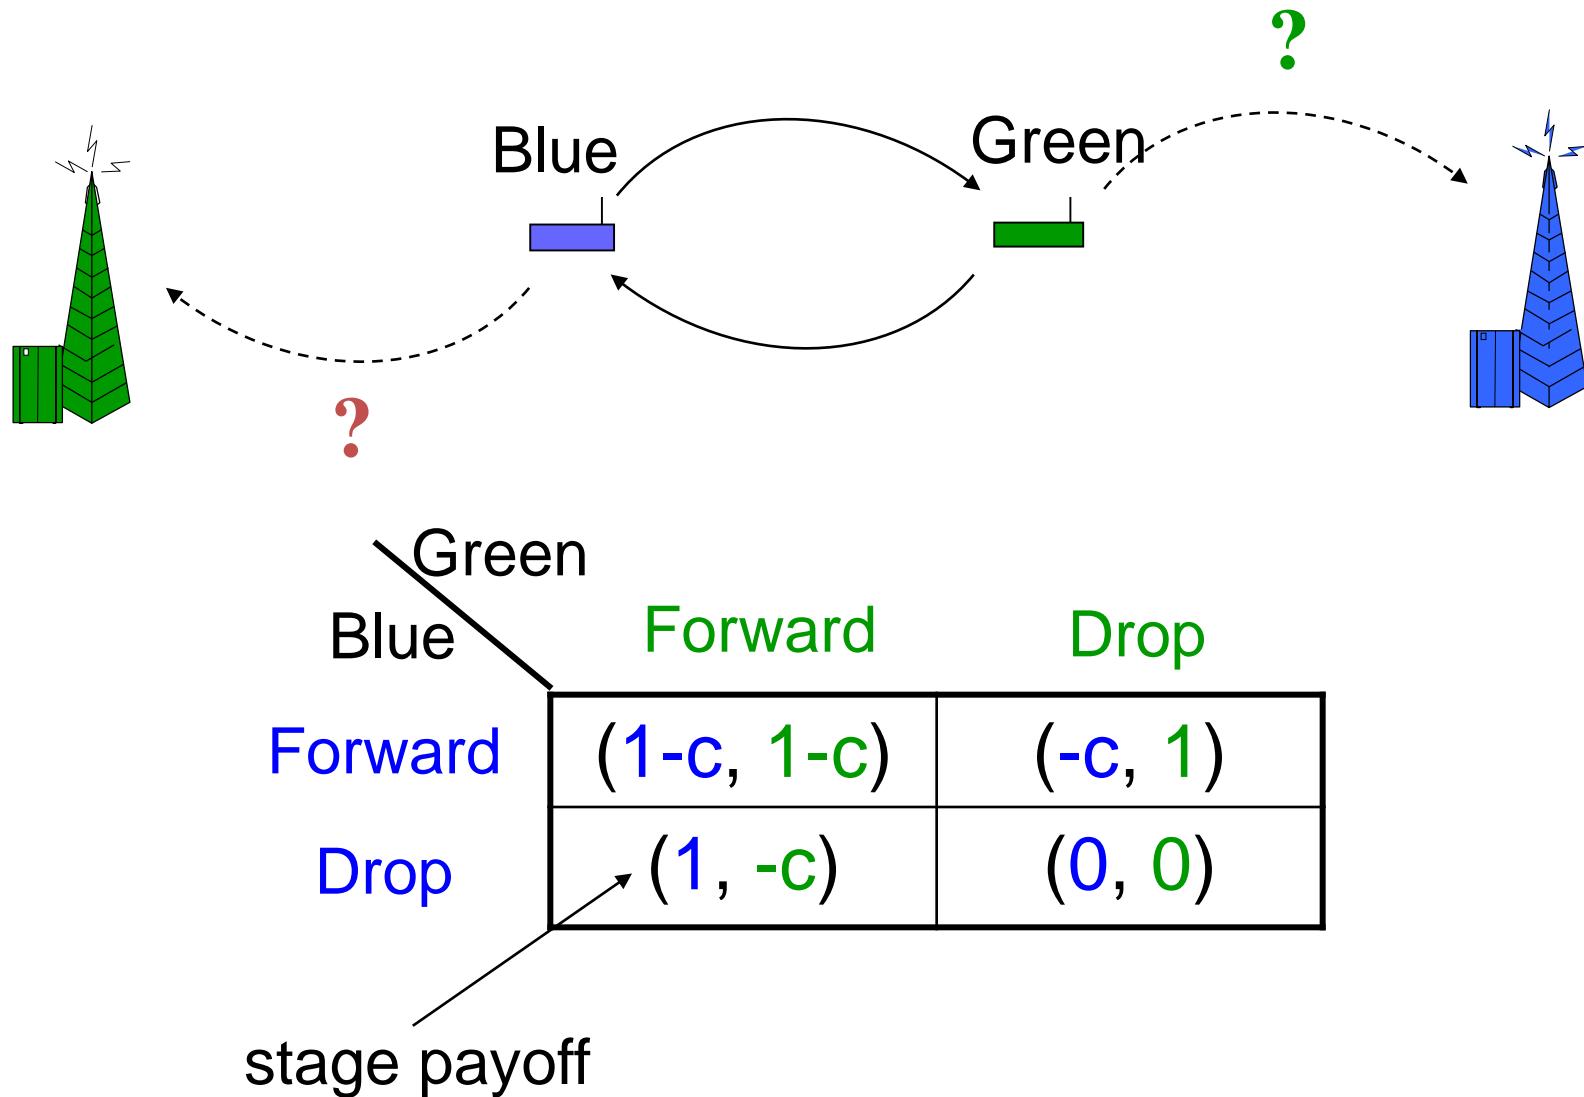
$$3-2 \stackrel{?}{\leq} [(2 \text{ “forever”}) - (\text{value of “0” for tomorrow and then “2” forever starting the next day})] \delta$$

$$3-2 \leq [2/(1-\delta) - 2\delta/(1-\delta)] \delta$$

$$\frac{1}{2} \leq \delta$$

Shorter punishments need more weight (δ) on future

The Repeated Forwarder's Dilemma



NE in Finite Repeated FD

In the finite-horizon Repeated Forwarder's Dilemma, the strategy profile (All-D, All-D) is a Nash equilibrium.

Payoffs in the Repeated Game FD

- Finite-horizon vs. infinite-horizon games
- Myopic vs. long-sighted repeated game

$$\text{myopic: } \bar{u}_i = u_i(t+1)$$

$$\text{long-sighted finite: } \bar{u}_i = \sum_{t=0}^T u_i(t)$$

$$\text{long-sighted infinite: } \bar{u}_i = \sum_{t=0}^{\infty} u_i(t)$$

$$\text{payoff with discounting: } \bar{u}_i = \sum_{t=0}^{\infty} u_i(t) \cdot \delta^t$$

$0 < \delta \leq 1$ is the discounting factor

Strategies in the Repeated Game FD

- usually, history-1 strategies, based on different inputs:
 - others' behavior: $m_i(t+1) = s_i[m_{-i}(t)]$
 - others' and own behavior: $m_i(t+1) = s_i[m_i(t), m_{-i}(t)]$
 - payoff: $m_i(t+1) = s_i[u_i(t)]$

Example strategies in the Forwarder's Dilemma:

Blue (t)	initial move	F	D	strategy name
Green (t+1)	F	F	F	AllC
	F	F	D	Tit-For-Tat (TFT)
	D	D	D	AllD
	F	D	F	Anti-TFT

Analysis of the Repeated Forwarder's Dilemma (I/3)

Infinite game with discounting:

$$\bar{u}_i = \sum_{t=0}^{\infty} u_i(t) \cdot \delta^t$$

Blue strategy	Green strategy
AIID	AIID
AIID	TFT
AIID	AIIC
AIIC	AIIC
AIIC	TFT
TFT	TFT

Blue utility	Green utility
0	0
1	-c
$1/(1-\delta)$	$-c/(1-\delta)$
$(1-c)/(1-\delta)$	$(1-c)/(1-\delta)$
$(1-c)/(1-\delta)$	$(1-c)/(1-\delta)$
$(1-c)/(1-\delta)$	$(1-c)/(1-\delta)$

Discount Factor Interpretation:

1. The player cares more about the near term payoff than in the long term payoff
2. The player has no preferences, but the game ends with probability of $1-\delta$ in each stage

Analysis of the Repeated Forwarder's Dilemma (2/3)

Blue strategy	Green strategy
AIID	AIID
AIID	TFT
AIID	AIIC
AIIC	AIIC
AIIC	TFT
TFT	TFT

Blue utility	Green utility
0	0
1	-c
$1/(1-\delta)$	$-c/(1-\delta)$
$(1-c)/(1-\delta)$	$(1-c)/(1-\delta)$
$(1-c)/(1-\delta)$	$(1-c)/(1-\delta)$
$(1-c)/(1-\delta)$	$(1-c)/(1-\delta)$

- AIIC receives a high payoff with itself and TFT, but
- AIID exploits AIIC
- AIID performs poor with itself
- TFT performs well with AIIC and itself, and
- TFT retaliates the defection of AIID

TFT is the best strategy if δ is high enough!

NE in Infinite Repeated FD

In the Repeated Forwarder's Dilemma, if both players play AIID, it is a Nash equilibrium.

NE in Infinite Repeated FD

In the Repeated Forwarder's Dilemma, both players playing TFT is a Nash equilibrium (if $\delta > c$).

Sketch of Proof:

If one deviate in stage t , then its payoff is:

$$(1-\delta) [(1+\delta+\delta^2 \dots +\delta^{t-1})(1-c) + \delta^t] = \\ 1-c+\delta^t(c-\delta) \rightarrow$$

Hence if “ $\delta > c$ ” there is no temptation to deviate

Or (i.e., other approach):

$$1-(1-c) \leq \delta \text{ (u(C,C) forever - u (D,D) forever)}$$

$$c \leq \delta ((1-c)/(1-\delta) - 0) \rightarrow \delta > c$$

Pareto-optimal in Repeated FD

The Nash equilibrium $s_{\text{Blue}} = \text{TFT}$ and $s_{\text{Green}} = \text{TFT}$ is Pareto-optimal
(but $s_{\text{Blue}} = \text{AllD}$ and $s_{\text{Green}} = \text{AllD}$ is not) !

Sketch of Proof:

There is no way for a player to go above his normalized payoff of $1-c$ without hurting his opponent's payoff

Formal Definition!

EQUILIBRIUM IN INFINITE REPEATED GAME

Minmax Value

The minmax value is the lowest stage payoff that the opponents of player i can force him to obtain with punishments, provided that i plays the best response against them.

$$\underline{u}_i = \min_{s_{-i}} \left[\max_{s_i} u_i(s_i, s_{-i}) \right]$$

Enforceable Payoff Profile

A payoff profile $u = (u_1, u_2, \dots, u_n)$ is
enforceable if $u_i \geq \underline{u}_i$

Feasible Payoff Profile

In n-player game $G=(N, S, u)$, a payoff profile u is **feasible** if there exist *fractional, non-negative values* α_j such that for all j , we can express u_i as $\sum_{j \in |S|} \alpha_j u_i(j)$, with $\sum_{j \in S} \alpha_j = 1$.

Example of Feasible Profile

(2,0)	(0,0)
(0,0)	(0,2)

- (1,1) is a feasible payoff given that we assign 0.5 to the strategy profile over diagonal.
- (2,2) is not feasible

Theorem

*Player i 's normalized payoff is at least equal to
Minmax value in any Nash equilibrium
of the infinitely repeated game,
regardless of the level of the discount factor*

Intuition: a player playing All-D will obtain a (normalized) payoff of at least 0

Folk Theorem

(Infinitely Repeated Game with Average Rewards)

Consider any n-player game G and any payoff vector (u_1, u_2, \dots, u_n) .

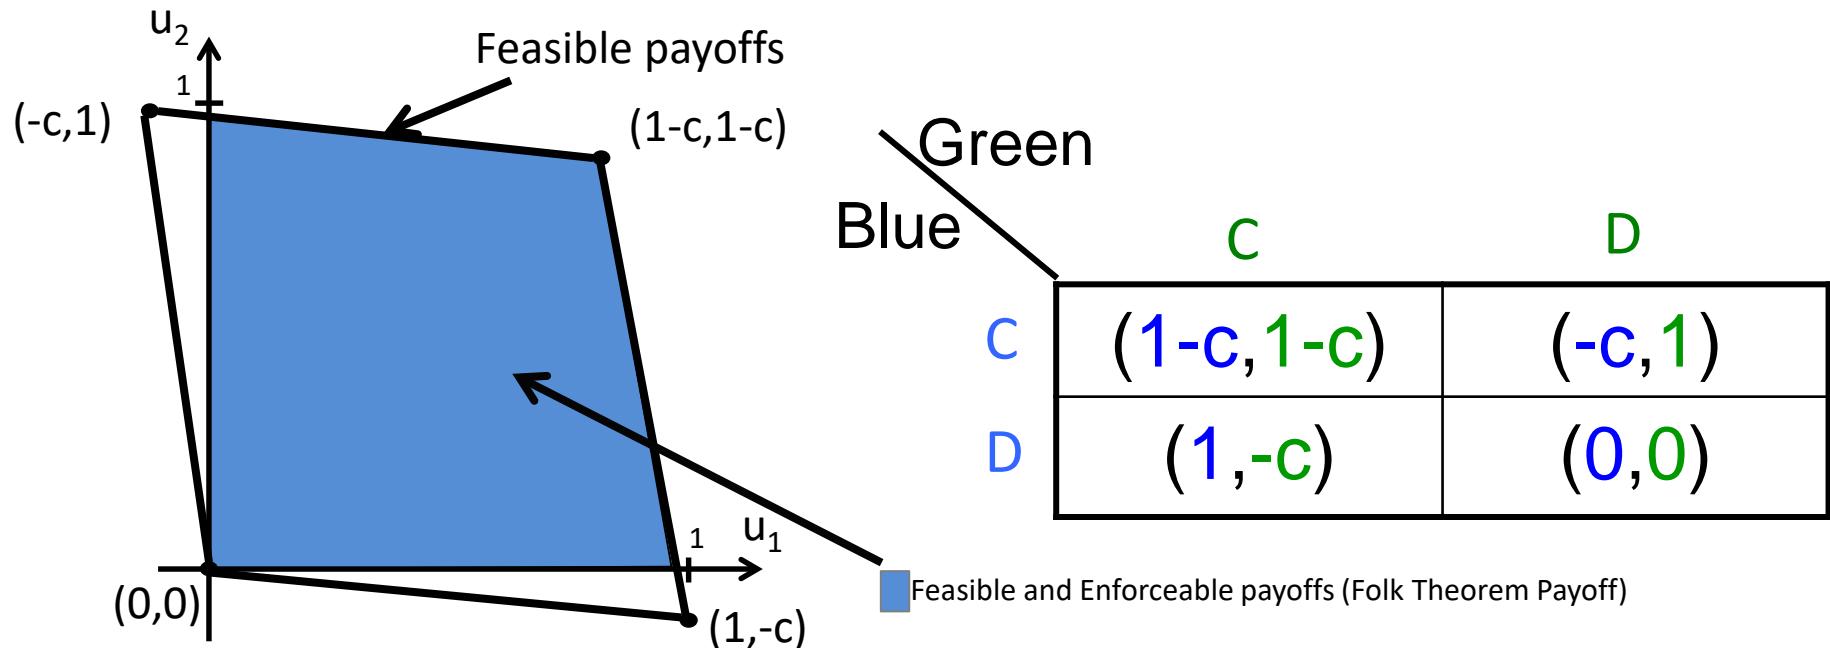
1. If u is the payoff in any Nash equilibrium of the infinitely repeated G with average rewards, then for each player i , u_i is enforceable.
2. If u is both feasible and enforceable, then u is the payoff in **some Nash equilibrium** of the infinitely repeated G with average rewards.

Folk Theorem

Infinitely Repeated Game with discounting Factor

For every feasible payoff vector $u = \{u_i\}_i$ with $u_i > \underline{u}_i$ (i.e., it is enforceable as well), there exists a discounting factor $\underline{\delta} < 1$, such that for all $\underline{\delta} < \delta < 1$, there is a Nash equilibrium with payoff u .

Feasible Payoffs in the Repeated Forwarder's Dilemma



Note that p_i can obtain at least his minmax value in any stage (*enforceable payoffs are always higher than the minmax payoff*)

Intuition and Example!

If the game is long enough, the gain obtained by a player by deviating once is outweighed by the loss in every subsequent period, when loss is due to the punishment (minmax) strategy of the other players.

Example: In infinite repeated FD, a player is deterred from deviating, because the short term gain obtained by the deviation (1 instead of $1-c$) is outweighed by the risk of being minmaxed (for example using the Trigger strategy) by the other player (provided that $c < \delta$).